

# final

March 26, 2023

```
[1]: #!pip install torch torchvision
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

## Prepare for Dataset

```
[2]: transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
           'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
           'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
           'castle', 'caterpillar', 'cattle', 'chair',
           'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
           'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
           'elephant', 'flatfish', 'forest', 'fox', 'girl',
           'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
           'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
           'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
           'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
```

```
'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
'wolf', 'woman', 'worm')
```


```
#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



dolphin pine\_tree plate kangaroo

### Choose a Device

```
[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

## Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
padding=1)
        self.bn1 = nn.BatchNorm2d(64)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(128)
        self.pool2 = nn.MaxPool2d(2, 2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(256)
        self.pool3 = nn.MaxPool2d(2, 2)
        self.relu3 = nn.ReLU()
        self.fc1 = nn.Linear(256 * 4 * 4, 1024)
        self.bn4 = nn.BatchNorm1d(1024)
        self.relu4 = nn.ReLU()
        self.fc2 = nn.Linear(1024, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.conv1(x)))
        x = self.pool2(self.relu2(self.conv2(x)))
        x = self.pool3(self.relu3(self.conv3(x)))
        x = x.view(-1, 256 * 4 * 4)
        x = self.relu4(self.fc1(x))
        x = self.fc2(x)
        return x

net = Net()      # Create the network instance.
net.to(device)  # Move the network parameters to the specified device.
```

```
[5]: Net(
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
```

```

    (relu1): ReLU()
    (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (relu2): ReLU()
    (conv3): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (relu3): ReLU()
    (fc1): Linear(in_features=4096, out_features=1024, bias=True)
    (bn4): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu4): ReLU()
    (fc2): Linear(in_features=1024, out_features=100, bias=True)
)

```

## Optimizer and Loss Function

```

[6]: # We use cross-entropy as loss function.
    loss_func = nn.CrossEntropyLoss()
    # We use stochastic gradient descent (SGD) as optimizer.
    opt = optim.SGD(net.parameters(), lr=0.0005, momentum=0.9)

```

## Training Procedure

```

[7]: import sys
    from tqdm.notebook import tqdm

    avg_losses = []    # Avg. losses.
    epochs = 7         # Total epochs.
    print_freq = 500   # Print frequency.

    for epoch in range(epochs): # Loop over the dataset multiple times.
        running_loss = 0.0      # Initialize running loss.
        for i, data in enumerate(tqdm(trainloader), 0):
            # Get the inputs.
            inputs, labels = data

            # Move the inputs to the specified device.
            inputs, labels = inputs.to(device), labels.to(device)

            # Zero the parameter gradients.
            opt.zero_grad()

```

```

    # Forward step.
    outputs = net(inputs)
    loss = loss_func(outputs, labels)

    # Backward step.
    loss.backward()

    # Optimization step (update the parameters).
    opt.step()

    # Print statistics.
    running_loss += loss.item()
    if i % print_freq == print_freq - 1: # Print every several mini-batches.
        avg_loss = running_loss / print_freq
        print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.
            ↪format(epoch, i, avg_loss), flush=True)
        sys.stdout.flush()
        avg_losses.append(avg_loss)
        running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 0, i:   499] avg mini-batch loss: 4.603
[epoch: 0, i:   999] avg mini-batch loss: 4.598
[epoch: 0, i:  1499] avg mini-batch loss: 4.593
[epoch: 0, i:  1999] avg mini-batch loss: 4.572
[epoch: 0, i:  2499] avg mini-batch loss: 4.537
[epoch: 0, i:  2999] avg mini-batch loss: 4.439
[epoch: 0, i:  3499] avg mini-batch loss: 4.342
[epoch: 0, i:  3999] avg mini-batch loss: 4.239
[epoch: 0, i:  4499] avg mini-batch loss: 4.162
[epoch: 0, i:  4999] avg mini-batch loss: 4.089
[epoch: 0, i:  5499] avg mini-batch loss: 4.036
[epoch: 0, i:  5999] avg mini-batch loss: 4.038
[epoch: 0, i:  6499] avg mini-batch loss: 3.987
[epoch: 0, i:  6999] avg mini-batch loss: 3.924
[epoch: 0, i:  7499] avg mini-batch loss: 3.885
[epoch: 0, i:  7999] avg mini-batch loss: 3.829
[epoch: 0, i:  8499] avg mini-batch loss: 3.855
[epoch: 0, i:  8999] avg mini-batch loss: 3.798
[epoch: 0, i:  9499] avg mini-batch loss: 3.743
[epoch: 0, i:  9999] avg mini-batch loss: 3.743
[epoch: 0, i: 10499] avg mini-batch loss: 3.712
[epoch: 0, i: 10999] avg mini-batch loss: 3.656
[epoch: 0, i: 11499] avg mini-batch loss: 3.637

```

[epoch: 0, i: 11999] avg mini-batch loss: 3.615  
[epoch: 0, i: 12499] avg mini-batch loss: 3.552

0%| | 0/12500 [00:00<?, ?it/s]

[epoch: 1, i: 499] avg mini-batch loss: 3.484  
[epoch: 1, i: 999] avg mini-batch loss: 3.514  
[epoch: 1, i: 1499] avg mini-batch loss: 3.426  
[epoch: 1, i: 1999] avg mini-batch loss: 3.432  
[epoch: 1, i: 2499] avg mini-batch loss: 3.363  
[epoch: 1, i: 2999] avg mini-batch loss: 3.370  
[epoch: 1, i: 3499] avg mini-batch loss: 3.290  
[epoch: 1, i: 3999] avg mini-batch loss: 3.296  
[epoch: 1, i: 4499] avg mini-batch loss: 3.317  
[epoch: 1, i: 4999] avg mini-batch loss: 3.288  
[epoch: 1, i: 5499] avg mini-batch loss: 3.293  
[epoch: 1, i: 5999] avg mini-batch loss: 3.263  
[epoch: 1, i: 6499] avg mini-batch loss: 3.175  
[epoch: 1, i: 6999] avg mini-batch loss: 3.236  
[epoch: 1, i: 7499] avg mini-batch loss: 3.203  
[epoch: 1, i: 7999] avg mini-batch loss: 3.098  
[epoch: 1, i: 8499] avg mini-batch loss: 3.130  
[epoch: 1, i: 8999] avg mini-batch loss: 3.045  
[epoch: 1, i: 9499] avg mini-batch loss: 3.118  
[epoch: 1, i: 9999] avg mini-batch loss: 3.048  
[epoch: 1, i: 10499] avg mini-batch loss: 3.156  
[epoch: 1, i: 10999] avg mini-batch loss: 3.025  
[epoch: 1, i: 11499] avg mini-batch loss: 2.982  
[epoch: 1, i: 11999] avg mini-batch loss: 3.031  
[epoch: 1, i: 12499] avg mini-batch loss: 2.957

0%| | 0/12500 [00:00<?, ?it/s]

[epoch: 2, i: 499] avg mini-batch loss: 2.813  
[epoch: 2, i: 999] avg mini-batch loss: 2.789  
[epoch: 2, i: 1499] avg mini-batch loss: 2.846  
[epoch: 2, i: 1999] avg mini-batch loss: 2.879  
[epoch: 2, i: 2499] avg mini-batch loss: 2.915  
[epoch: 2, i: 2999] avg mini-batch loss: 2.844  
[epoch: 2, i: 3499] avg mini-batch loss: 2.730  
[epoch: 2, i: 3999] avg mini-batch loss: 2.768  
[epoch: 2, i: 4499] avg mini-batch loss: 2.765  
[epoch: 2, i: 4999] avg mini-batch loss: 2.822  
[epoch: 2, i: 5499] avg mini-batch loss: 2.750  
[epoch: 2, i: 5999] avg mini-batch loss: 2.754  
[epoch: 2, i: 6499] avg mini-batch loss: 2.749  
[epoch: 2, i: 6999] avg mini-batch loss: 2.687  
[epoch: 2, i: 7499] avg mini-batch loss: 2.730  
[epoch: 2, i: 7999] avg mini-batch loss: 2.721  
[epoch: 2, i: 8499] avg mini-batch loss: 2.701

```

[epoch: 2, i: 8999] avg mini-batch loss: 2.690
[epoch: 2, i: 9499] avg mini-batch loss: 2.621
[epoch: 2, i: 9999] avg mini-batch loss: 2.630
[epoch: 2, i: 10499] avg mini-batch loss: 2.651
[epoch: 2, i: 10999] avg mini-batch loss: 2.647
[epoch: 2, i: 11499] avg mini-batch loss: 2.603
[epoch: 2, i: 11999] avg mini-batch loss: 2.602
[epoch: 2, i: 12499] avg mini-batch loss: 2.593

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 3, i: 499] avg mini-batch loss: 2.361
[epoch: 3, i: 999] avg mini-batch loss: 2.410
[epoch: 3, i: 1499] avg mini-batch loss: 2.350
[epoch: 3, i: 1999] avg mini-batch loss: 2.399
[epoch: 3, i: 2499] avg mini-batch loss: 2.407
[epoch: 3, i: 2999] avg mini-batch loss: 2.373
[epoch: 3, i: 3499] avg mini-batch loss: 2.350
[epoch: 3, i: 3999] avg mini-batch loss: 2.412
[epoch: 3, i: 4499] avg mini-batch loss: 2.423
[epoch: 3, i: 4999] avg mini-batch loss: 2.423
[epoch: 3, i: 5499] avg mini-batch loss: 2.445
[epoch: 3, i: 5999] avg mini-batch loss: 2.323
[epoch: 3, i: 6499] avg mini-batch loss: 2.386
[epoch: 3, i: 6999] avg mini-batch loss: 2.291
[epoch: 3, i: 7499] avg mini-batch loss: 2.346
[epoch: 3, i: 7999] avg mini-batch loss: 2.348
[epoch: 3, i: 8499] avg mini-batch loss: 2.319
[epoch: 3, i: 8999] avg mini-batch loss: 2.345
[epoch: 3, i: 9499] avg mini-batch loss: 2.314
[epoch: 3, i: 9999] avg mini-batch loss: 2.328
[epoch: 3, i: 10499] avg mini-batch loss: 2.349
[epoch: 3, i: 10999] avg mini-batch loss: 2.334
[epoch: 3, i: 11499] avg mini-batch loss: 2.247
[epoch: 3, i: 11999] avg mini-batch loss: 2.334
[epoch: 3, i: 12499] avg mini-batch loss: 2.304

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 4, i: 499] avg mini-batch loss: 1.921
[epoch: 4, i: 999] avg mini-batch loss: 2.030
[epoch: 4, i: 1499] avg mini-batch loss: 2.007
[epoch: 4, i: 1999] avg mini-batch loss: 2.024
[epoch: 4, i: 2499] avg mini-batch loss: 2.050
[epoch: 4, i: 2999] avg mini-batch loss: 1.966
[epoch: 4, i: 3499] avg mini-batch loss: 2.048
[epoch: 4, i: 3999] avg mini-batch loss: 2.025
[epoch: 4, i: 4499] avg mini-batch loss: 2.008
[epoch: 4, i: 4999] avg mini-batch loss: 2.039
[epoch: 4, i: 5499] avg mini-batch loss: 2.038

```

```

[epoch: 4, i: 5999] avg mini-batch loss: 2.092
[epoch: 4, i: 6499] avg mini-batch loss: 2.013
[epoch: 4, i: 6999] avg mini-batch loss: 2.021
[epoch: 4, i: 7499] avg mini-batch loss: 2.080
[epoch: 4, i: 7999] avg mini-batch loss: 2.009
[epoch: 4, i: 8499] avg mini-batch loss: 2.073
[epoch: 4, i: 8999] avg mini-batch loss: 1.984
[epoch: 4, i: 9499] avg mini-batch loss: 1.955
[epoch: 4, i: 9999] avg mini-batch loss: 1.977
[epoch: 4, i: 10499] avg mini-batch loss: 1.967
[epoch: 4, i: 10999] avg mini-batch loss: 1.970
[epoch: 4, i: 11499] avg mini-batch loss: 1.981
[epoch: 4, i: 11999] avg mini-batch loss: 2.001
[epoch: 4, i: 12499] avg mini-batch loss: 2.021

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 5, i: 499] avg mini-batch loss: 1.577
[epoch: 5, i: 999] avg mini-batch loss: 1.627
[epoch: 5, i: 1499] avg mini-batch loss: 1.617
[epoch: 5, i: 1999] avg mini-batch loss: 1.618
[epoch: 5, i: 2499] avg mini-batch loss: 1.671
[epoch: 5, i: 2999] avg mini-batch loss: 1.634
[epoch: 5, i: 3499] avg mini-batch loss: 1.651
[epoch: 5, i: 3999] avg mini-batch loss: 1.624
[epoch: 5, i: 4499] avg mini-batch loss: 1.712
[epoch: 5, i: 4999] avg mini-batch loss: 1.606
[epoch: 5, i: 5499] avg mini-batch loss: 1.683
[epoch: 5, i: 5999] avg mini-batch loss: 1.664
[epoch: 5, i: 6499] avg mini-batch loss: 1.650
[epoch: 5, i: 6999] avg mini-batch loss: 1.688
[epoch: 5, i: 7499] avg mini-batch loss: 1.733
[epoch: 5, i: 7999] avg mini-batch loss: 1.671
[epoch: 5, i: 8499] avg mini-batch loss: 1.699
[epoch: 5, i: 8999] avg mini-batch loss: 1.688
[epoch: 5, i: 9499] avg mini-batch loss: 1.779
[epoch: 5, i: 9999] avg mini-batch loss: 1.662
[epoch: 5, i: 10499] avg mini-batch loss: 1.711
[epoch: 5, i: 10999] avg mini-batch loss: 1.745
[epoch: 5, i: 11499] avg mini-batch loss: 1.666
[epoch: 5, i: 11999] avg mini-batch loss: 1.647
[epoch: 5, i: 12499] avg mini-batch loss: 1.742

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 6, i: 499] avg mini-batch loss: 1.158
[epoch: 6, i: 999] avg mini-batch loss: 1.198
[epoch: 6, i: 1499] avg mini-batch loss: 1.205
[epoch: 6, i: 1999] avg mini-batch loss: 1.234
[epoch: 6, i: 2499] avg mini-batch loss: 1.223

```



```

[epoch: 6, i: 2999] avg mini-batch loss: 1.223
[epoch: 6, i: 3499] avg mini-batch loss: 1.259
[epoch: 6, i: 3999] avg mini-batch loss: 1.239
[epoch: 6, i: 4499] avg mini-batch loss: 1.283
[epoch: 6, i: 4999] avg mini-batch loss: 1.283
[epoch: 6, i: 5499] avg mini-batch loss: 1.346
[epoch: 6, i: 5999] avg mini-batch loss: 1.281
[epoch: 6, i: 6499] avg mini-batch loss: 1.327
[epoch: 6, i: 6999] avg mini-batch loss: 1.335
[epoch: 6, i: 7499] avg mini-batch loss: 1.315
[epoch: 6, i: 7999] avg mini-batch loss: 1.360
[epoch: 6, i: 8499] avg mini-batch loss: 1.345
[epoch: 6, i: 8999] avg mini-batch loss: 1.451
[epoch: 6, i: 9499] avg mini-batch loss: 1.354
[epoch: 6, i: 9999] avg mini-batch loss: 1.346
[epoch: 6, i: 10499] avg mini-batch loss: 1.345
[epoch: 6, i: 10999] avg mini-batch loss: 1.352
[epoch: 6, i: 11499] avg mini-batch loss: 1.382
[epoch: 6, i: 11999] avg mini-batch loss: 1.431
[epoch: 6, i: 12499] avg mini-batch loss: 1.424
Finished Training.


```

### Training Loss Curve

```

[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()


```



## Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                              for j in range(4)))
```



GroundTruth: mountain forest seal mushroom

Predicted: spider forest camel trout

```
[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 42 %

```
[11]: # Get test accuracy for each class.
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

```
-----
IndexError                                Traceback (most recent call last)
/tmp/ipykernel_6876/2959918759.py in <module>
    11         for i in range(4):
    12             label = labels[i]
----> 13             class_correct[label] += c[i].item()
    14             class_total[label] += 1
    15

IndexError: list index out of range
```

```
[ ]: # One of the changes I made was that I added another layer in the network
# that takes the output from the second convolutional layer, applies ReLU
    ↪ activation,
# and passes it through a 2x2 AvgPool layer to capture more relationships.
# I also reduced the learning rate of the optimizer to half (0.0005) which
# which allowed the optimizer to take smaller steps towards the minimum of
# the loss function which might allow for a better chance of finding the global
    ↪ min of loss.
```