

# final\_8

March 26, 2023

```
[1]: #!pip install torch torchvision
      %matplotlib inline
      import matplotlib.pyplot as plt
      import numpy as np
      import torch
      import torchvision
      import torchvision.transforms as transforms
      import torch.nn as nn
      import torch.nn.functional as F
      import torch.optim as optim
```

## Prepare for Dataset

```
[2]: transform = transforms.Compose(
      [transforms.ToTensor(),
       transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

      trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                                download=True, transform=transform)
      trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                                shuffle=True, num_workers=2)

      testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                                download=True, transform=transform)
      testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                                shuffle=False, num_workers=2)

      classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
                  'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
                  'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
                  'castle', 'caterpillar', 'cattle', 'chair',
                  'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
                  'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
                  'elephant', 'flatfish', 'forest', 'fox', 'girl',
                  'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
                  'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
                  'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
                  'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
```

```

    'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
    'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
    'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
    'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
    'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
    'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
    'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
    'wolf', 'woman', 'worm')

#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```


Files already downloaded and verified  
Files already downloaded and verified

```

[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))

```



lizard whale lizard lawn\_mower

### Choose a Device

```

[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

```

```
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

## Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
padding=1)
        self.bn1 = nn.BatchNorm2d(64)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(128)
        self.pool2 = nn.MaxPool2d(2, 2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(256)
        self.pool3 = nn.MaxPool2d(2, 2)
        self.relu3 = nn.SELU()
        self.fc1 = nn.Linear(256 * 4 * 4, 1024)
        self.bn4 = nn.BatchNorm1d(1024)
        self.relu4 = nn.SELU()
        self.fc2 = nn.Linear(1024, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.conv1(x)))
        x = self.pool2(self.relu2(self.conv2(x)))
        x = self.pool3(self.relu3(self.conv3(x)))
        x = x.view(-1, 256 * 4 * 4)
        x = self.relu4(self.fc1(x))
        x = self.fc2(x)
        return x

net = Net()      # Create the network instance.
net.to(device)  # Move the network parameters to the specified device.
```

```
[5]: Net(
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
```

```

    (relu1): ReLU()
    (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (relu2): ReLU()
    (conv3): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (relu3): SELU()
    (fc1): Linear(in_features=4096, out_features=1024, bias=True)
    (bn4): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu4): SELU()
    (fc2): Linear(in_features=1024, out_features=100, bias=True)
)

```

## Optimizer and Loss Function

```

[6]: # We use cross-entropy as loss function.
    loss_func = nn.CrossEntropyLoss()
    # We use stochastic gradient descent (SGD) as optimizer.
    #opt = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)
    opt = optim.Adam(net.parameters(), lr=0.0005)

```

## Training Procedure

```

[7]: import sys
    from tqdm.notebook import tqdm

    avg_losses = []    # Avg. losses.
    epochs = 8         # Total epochs.
    print_freq = 500   # Print frequency.

    for epoch in range(epochs): # Loop over the dataset multiple times.
        running_loss = 0.0    # Initialize running loss.
        for i, data in enumerate(tqdm(trainloader), 0):
            # Get the inputs.
            inputs, labels = data

            # Move the inputs to the specified device.
            inputs, labels = inputs.to(device), labels.to(device)

            # Zero the parameter gradients.

```

```

opt.zero_grad()

# Forward step.
outputs = net(inputs)
loss = loss_func(outputs, labels)

# Backward step.
loss.backward()

# Optimization step (update the parameters).
opt.step()

# Print statistics.
running_loss += loss.item()
if i % print_freq == print_freq - 1: # Print every several mini-batches.
    avg_loss = running_loss / print_freq
    print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.
    ↪format(epoch, i, avg_loss),flush=True)
    sys.stdout.flush()
    avg_losses.append(avg_loss)
    running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 0, i:   499] avg mini-batch loss: 4.433
[epoch: 0, i:   999] avg mini-batch loss: 4.165
[epoch: 0, i:  1499] avg mini-batch loss: 4.040
[epoch: 0, i:  1999] avg mini-batch loss: 3.907
[epoch: 0, i:  2499] avg mini-batch loss: 3.804
[epoch: 0, i:  2999] avg mini-batch loss: 3.753
[epoch: 0, i:  3499] avg mini-batch loss: 3.644
[epoch: 0, i:  3999] avg mini-batch loss: 3.572
[epoch: 0, i:  4499] avg mini-batch loss: 3.536
[epoch: 0, i:  4999] avg mini-batch loss: 3.485
[epoch: 0, i:  5499] avg mini-batch loss: 3.422
[epoch: 0, i:  5999] avg mini-batch loss: 3.298
[epoch: 0, i:  6499] avg mini-batch loss: 3.325
[epoch: 0, i:  6999] avg mini-batch loss: 3.312
[epoch: 0, i:  7499] avg mini-batch loss: 3.102
[epoch: 0, i:  7999] avg mini-batch loss: 3.241
[epoch: 0, i:  8499] avg mini-batch loss: 3.180
[epoch: 0, i:  8999] avg mini-batch loss: 3.127
[epoch: 0, i:  9499] avg mini-batch loss: 3.099
[epoch: 0, i:  9999] avg mini-batch loss: 2.987
[epoch: 0, i: 10499] avg mini-batch loss: 3.111
[epoch: 0, i: 10999] avg mini-batch loss: 3.056

```

```
[epoch: 0, i: 11499] avg mini-batch loss: 3.053
[epoch: 0, i: 11999] avg mini-batch loss: 3.102
[epoch: 0, i: 12499] avg mini-batch loss: 3.026
```

```
0%|          | 0/12500 [00:00<?, ?it/s]
```

```
[epoch: 1, i: 499] avg mini-batch loss: 2.803
[epoch: 1, i: 999] avg mini-batch loss: 2.857
[epoch: 1, i: 1499] avg mini-batch loss: 2.861
[epoch: 1, i: 1999] avg mini-batch loss: 2.862
[epoch: 1, i: 2499] avg mini-batch loss: 2.911
[epoch: 1, i: 2999] avg mini-batch loss: 2.805
[epoch: 1, i: 3499] avg mini-batch loss: 2.863
[epoch: 1, i: 3999] avg mini-batch loss: 2.870
[epoch: 1, i: 4499] avg mini-batch loss: 2.809
[epoch: 1, i: 4999] avg mini-batch loss: 2.827
[epoch: 1, i: 5499] avg mini-batch loss: 2.861
[epoch: 1, i: 5999] avg mini-batch loss: 2.847
[epoch: 1, i: 6499] avg mini-batch loss: 2.856
[epoch: 1, i: 6999] avg mini-batch loss: 2.783
[epoch: 1, i: 7499] avg mini-batch loss: 2.864
[epoch: 1, i: 7999] avg mini-batch loss: 2.801
[epoch: 1, i: 8499] avg mini-batch loss: 2.800
[epoch: 1, i: 8999] avg mini-batch loss: 2.753
[epoch: 1, i: 9499] avg mini-batch loss: 2.796
[epoch: 1, i: 9999] avg mini-batch loss: 2.774
[epoch: 1, i: 10499] avg mini-batch loss: 2.789
[epoch: 1, i: 10999] avg mini-batch loss: 2.733
[epoch: 1, i: 11499] avg mini-batch loss: 2.828
[epoch: 1, i: 11999] avg mini-batch loss: 2.741
[epoch: 1, i: 12499] avg mini-batch loss: 2.749
```

```
0%|          | 0/12500 [00:00<?, ?it/s]
```

```
[epoch: 2, i: 499] avg mini-batch loss: 2.506
[epoch: 2, i: 999] avg mini-batch loss: 2.439
[epoch: 2, i: 1499] avg mini-batch loss: 2.604
[epoch: 2, i: 1999] avg mini-batch loss: 2.493
[epoch: 2, i: 2499] avg mini-batch loss: 2.519
[epoch: 2, i: 2999] avg mini-batch loss: 2.581
[epoch: 2, i: 3499] avg mini-batch loss: 2.569
[epoch: 2, i: 3999] avg mini-batch loss: 2.652
[epoch: 2, i: 4499] avg mini-batch loss: 2.554
[epoch: 2, i: 4999] avg mini-batch loss: 2.548
[epoch: 2, i: 5499] avg mini-batch loss: 2.626
[epoch: 2, i: 5999] avg mini-batch loss: 2.638
[epoch: 2, i: 6499] avg mini-batch loss: 2.550
[epoch: 2, i: 6999] avg mini-batch loss: 2.524
[epoch: 2, i: 7499] avg mini-batch loss: 2.557
[epoch: 2, i: 7999] avg mini-batch loss: 2.592
```

```

[epoch: 2, i: 8499] avg mini-batch loss: 2.567
[epoch: 2, i: 8999] avg mini-batch loss: 2.573
[epoch: 2, i: 9499] avg mini-batch loss: 2.581
[epoch: 2, i: 9999] avg mini-batch loss: 2.564
[epoch: 2, i: 10499] avg mini-batch loss: 2.607
[epoch: 2, i: 10999] avg mini-batch loss: 2.569
[epoch: 2, i: 11499] avg mini-batch loss: 2.572
[epoch: 2, i: 11999] avg mini-batch loss: 2.527
[epoch: 2, i: 12499] avg mini-batch loss: 2.602

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 3, i: 499] avg mini-batch loss: 2.155
[epoch: 3, i: 999] avg mini-batch loss: 2.164
[epoch: 3, i: 1499] avg mini-batch loss: 2.245
[epoch: 3, i: 1999] avg mini-batch loss: 2.269
[epoch: 3, i: 2499] avg mini-batch loss: 2.285
[epoch: 3, i: 2999] avg mini-batch loss: 2.293
[epoch: 3, i: 3499] avg mini-batch loss: 2.305
[epoch: 3, i: 3999] avg mini-batch loss: 2.323
[epoch: 3, i: 4499] avg mini-batch loss: 2.388
[epoch: 3, i: 4999] avg mini-batch loss: 2.375
[epoch: 3, i: 5499] avg mini-batch loss: 2.366
[epoch: 3, i: 5999] avg mini-batch loss: 2.388
[epoch: 3, i: 6499] avg mini-batch loss: 2.343
[epoch: 3, i: 6999] avg mini-batch loss: 2.311
[epoch: 3, i: 7499] avg mini-batch loss: 2.421
[epoch: 3, i: 7999] avg mini-batch loss: 2.413
[epoch: 3, i: 8499] avg mini-batch loss: 2.346
[epoch: 3, i: 8999] avg mini-batch loss: 2.446
[epoch: 3, i: 9499] avg mini-batch loss: 2.358
[epoch: 3, i: 9999] avg mini-batch loss: 2.368
[epoch: 3, i: 10499] avg mini-batch loss: 2.454
[epoch: 3, i: 10999] avg mini-batch loss: 2.408
[epoch: 3, i: 11499] avg mini-batch loss: 2.412
[epoch: 3, i: 11999] avg mini-batch loss: 2.420
[epoch: 3, i: 12499] avg mini-batch loss: 2.490

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 4, i: 499] avg mini-batch loss: 1.873
[epoch: 4, i: 999] avg mini-batch loss: 1.978
[epoch: 4, i: 1499] avg mini-batch loss: 2.013
[epoch: 4, i: 1999] avg mini-batch loss: 2.012
[epoch: 4, i: 2499] avg mini-batch loss: 1.955
[epoch: 4, i: 2999] avg mini-batch loss: 2.038
[epoch: 4, i: 3499] avg mini-batch loss: 2.053
[epoch: 4, i: 3999] avg mini-batch loss: 2.058
[epoch: 4, i: 4499] avg mini-batch loss: 2.111
[epoch: 4, i: 4999] avg mini-batch loss: 2.133

```

```

[epoch: 4, i: 5499] avg mini-batch loss: 2.247
[epoch: 4, i: 5999] avg mini-batch loss: 2.190
[epoch: 4, i: 6499] avg mini-batch loss: 2.146
[epoch: 4, i: 6999] avg mini-batch loss: 2.061
[epoch: 4, i: 7499] avg mini-batch loss: 2.133
[epoch: 4, i: 7999] avg mini-batch loss: 2.162
[epoch: 4, i: 8499] avg mini-batch loss: 2.152
[epoch: 4, i: 8999] avg mini-batch loss: 2.160
[epoch: 4, i: 9499] avg mini-batch loss: 2.178
[epoch: 4, i: 9999] avg mini-batch loss: 2.247
[epoch: 4, i: 10499] avg mini-batch loss: 2.207
[epoch: 4, i: 10999] avg mini-batch loss: 2.208
[epoch: 4, i: 11499] avg mini-batch loss: 2.209
[epoch: 4, i: 11999] avg mini-batch loss: 2.292
[epoch: 4, i: 12499] avg mini-batch loss: 2.260

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 5, i: 499] avg mini-batch loss: 1.669
[epoch: 5, i: 999] avg mini-batch loss: 1.685
[epoch: 5, i: 1499] avg mini-batch loss: 1.691
[epoch: 5, i: 1999] avg mini-batch loss: 1.760
[epoch: 5, i: 2499] avg mini-batch loss: 1.895
[epoch: 5, i: 2999] avg mini-batch loss: 1.893
[epoch: 5, i: 3499] avg mini-batch loss: 1.765
[epoch: 5, i: 3999] avg mini-batch loss: 1.862
[epoch: 5, i: 4499] avg mini-batch loss: 1.862
[epoch: 5, i: 4999] avg mini-batch loss: 1.903
[epoch: 5, i: 5499] avg mini-batch loss: 1.867
[epoch: 5, i: 5999] avg mini-batch loss: 1.899
[epoch: 5, i: 6499] avg mini-batch loss: 1.901
[epoch: 5, i: 6999] avg mini-batch loss: 1.943
[epoch: 5, i: 7499] avg mini-batch loss: 1.977
[epoch: 5, i: 7999] avg mini-batch loss: 2.069
[epoch: 5, i: 8499] avg mini-batch loss: 1.950
[epoch: 5, i: 8999] avg mini-batch loss: 2.017
[epoch: 5, i: 9499] avg mini-batch loss: 1.961
[epoch: 5, i: 9999] avg mini-batch loss: 2.034
[epoch: 5, i: 10499] avg mini-batch loss: 1.942
[epoch: 5, i: 10999] avg mini-batch loss: 2.015
[epoch: 5, i: 11499] avg mini-batch loss: 2.075
[epoch: 5, i: 11999] avg mini-batch loss: 2.100
[epoch: 5, i: 12499] avg mini-batch loss: 2.000

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 6, i: 499] avg mini-batch loss: 1.400
[epoch: 6, i: 999] avg mini-batch loss: 1.474
[epoch: 6, i: 1499] avg mini-batch loss: 1.408
[epoch: 6, i: 1999] avg mini-batch loss: 1.543

```



```

[epoch: 6, i: 2499] avg mini-batch loss: 1.576
[epoch: 6, i: 2999] avg mini-batch loss: 1.543
[epoch: 6, i: 3499] avg mini-batch loss: 1.679
[epoch: 6, i: 3999] avg mini-batch loss: 1.648
[epoch: 6, i: 4499] avg mini-batch loss: 1.709
[epoch: 6, i: 4999] avg mini-batch loss: 1.650
[epoch: 6, i: 5499] avg mini-batch loss: 1.682
[epoch: 6, i: 5999] avg mini-batch loss: 1.674
[epoch: 6, i: 6499] avg mini-batch loss: 1.761
[epoch: 6, i: 6999] avg mini-batch loss: 1.741
[epoch: 6, i: 7499] avg mini-batch loss: 1.795
[epoch: 6, i: 7999] avg mini-batch loss: 1.776
[epoch: 6, i: 8499] avg mini-batch loss: 1.798
[epoch: 6, i: 8999] avg mini-batch loss: 1.790
[epoch: 6, i: 9499] avg mini-batch loss: 1.809
[epoch: 6, i: 9999] avg mini-batch loss: 1.847
[epoch: 6, i: 10499] avg mini-batch loss: 1.881
[epoch: 6, i: 10999] avg mini-batch loss: 1.801
[epoch: 6, i: 11499] avg mini-batch loss: 1.922
[epoch: 6, i: 11999] avg mini-batch loss: 1.887
[epoch: 6, i: 12499] avg mini-batch loss: 1.878

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

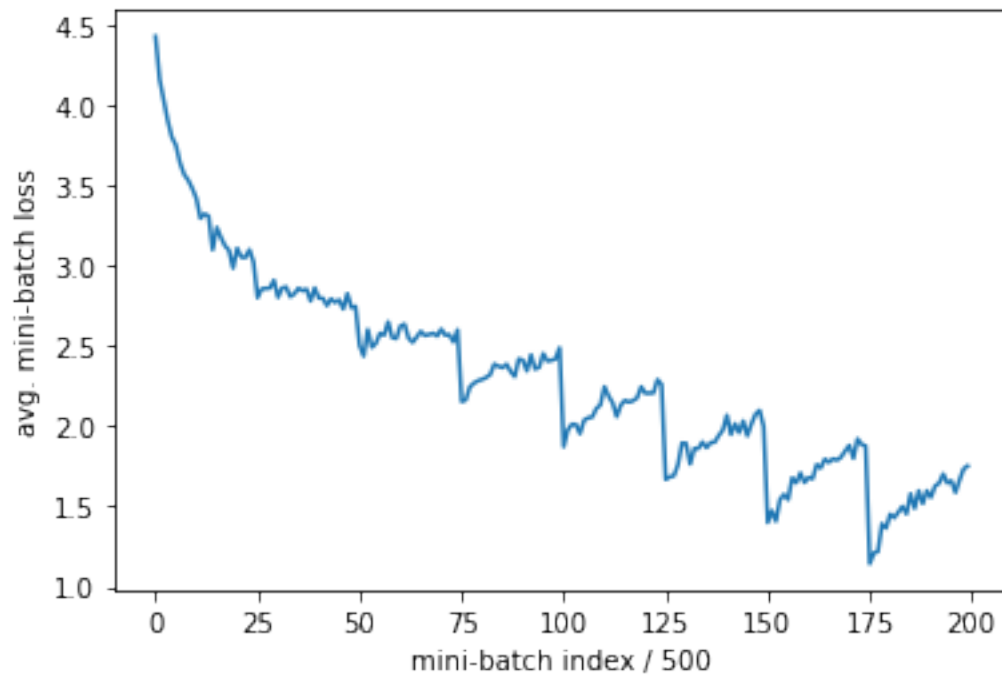
[epoch: 7, i: 499] avg mini-batch loss: 1.143
[epoch: 7, i: 999] avg mini-batch loss: 1.210
[epoch: 7, i: 1499] avg mini-batch loss: 1.219
[epoch: 7, i: 1999] avg mini-batch loss: 1.393
[epoch: 7, i: 2499] avg mini-batch loss: 1.365
[epoch: 7, i: 2999] avg mini-batch loss: 1.452
[epoch: 7, i: 3499] avg mini-batch loss: 1.430
[epoch: 7, i: 3999] avg mini-batch loss: 1.466
[epoch: 7, i: 4499] avg mini-batch loss: 1.501
[epoch: 7, i: 4999] avg mini-batch loss: 1.455
[epoch: 7, i: 5499] avg mini-batch loss: 1.580
[epoch: 7, i: 5999] avg mini-batch loss: 1.489
[epoch: 7, i: 6499] avg mini-batch loss: 1.598
[epoch: 7, i: 6999] avg mini-batch loss: 1.519
[epoch: 7, i: 7499] avg mini-batch loss: 1.599
[epoch: 7, i: 7999] avg mini-batch loss: 1.558
[epoch: 7, i: 8499] avg mini-batch loss: 1.628
[epoch: 7, i: 8999] avg mini-batch loss: 1.646
[epoch: 7, i: 9499] avg mini-batch loss: 1.703
[epoch: 7, i: 9999] avg mini-batch loss: 1.647
[epoch: 7, i: 10499] avg mini-batch loss: 1.660
[epoch: 7, i: 10999] avg mini-batch loss: 1.588
[epoch: 7, i: 11499] avg mini-batch loss: 1.666
[epoch: 7, i: 11999] avg mini-batch loss: 1.735
[epoch: 7, i: 12499] avg mini-batch loss: 1.753

```

Finished Training.

### Training Loss Curve


```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



### Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                               for j in range(4)))
```



GroundTruth: mountain forest seal mushroom  
 Predicted: dolphin squirrel seal seal

```
[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 30 %

```
[11]: # Get test accuracy for each class.
class_correct = [0] * len(classes)
class_total = [0] * len(classes)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(len(labels)):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1
```

```
for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

Accuracy of apple : 68 %  
Accuracy of aquarium\_fish : 30 %  
Accuracy of baby : 25 %  
Accuracy of bear : 9 %  
Accuracy of beaver : 19 %  
Accuracy of bed : 45 %  
Accuracy of bee : 35 %  
Accuracy of beetle : 22 %  
Accuracy of bicycle : 34 %  
Accuracy of bottle : 42 %  
Accuracy of bowl : 18 %  
Accuracy of boy : 15 %  
Accuracy of bridge : 24 %  
Accuracy of bus : 18 %  
Accuracy of butterfly : 22 %  
Accuracy of camel : 14 %  
Accuracy of can : 37 %  
Accuracy of castle : 29 %  
Accuracy of caterpillar : 19 %  
Accuracy of cattle : 13 %  
Accuracy of chair : 57 %  
Accuracy of chimpanzee : 48 %  
Accuracy of clock : 30 %  
Accuracy of cloud : 49 %  
Accuracy of cockroach : 35 %  
Accuracy of couch : 10 %  
Accuracy of crab : 27 %  
Accuracy of crocodile : 14 %  
Accuracy of cup : 43 %  
Accuracy of dinosaur : 21 %  
Accuracy of dolphin : 28 %  
Accuracy of elephant : 17 %  
Accuracy of flatfish : 44 %  
Accuracy of forest : 29 %  
Accuracy of fox : 19 %  
Accuracy of girl : 16 %  
Accuracy of hamster : 48 %  
Accuracy of house : 17 %  
Accuracy of kangaroo : 18 %  
Accuracy of keyboard : 22 %  
Accuracy of lamp : 31 %  
Accuracy of lawn\_mower : 63 %  
Accuracy of leopard : 25 %

Accuracy of lion : 13 %  
Accuracy of lizard : 13 %  
Accuracy of lobster : 13 %  
Accuracy of man : 19 %  
Accuracy of maple\_tree : 33 %  
Accuracy of motorcycle : 70 %  
Accuracy of mountain : 31 %  
Accuracy of mouse : 1 %  
Accuracy of mushroom : 19 %  
Accuracy of oak\_tree : 34 %  
Accuracy of orange : 56 %  
Accuracy of orchid : 27 %  
Accuracy of otter : 7 %  
Accuracy of palm\_tree : 34 %  
Accuracy of pear : 39 %  
Accuracy of pickup\_truck : 39 %  
Accuracy of pine\_tree : 43 %  
Accuracy of plain : 54 %  
Accuracy of plate : 33 %  
Accuracy of poppy : 43 %  
Accuracy of porcupine : 21 %  
Accuracy of possum : 11 %  
Accuracy of rabbit : 10 %  
Accuracy of raccoon : 13 %  
Accuracy of ray : 31 %  
Accuracy of road : 64 %  
Accuracy of rocket : 66 %  
Accuracy of rose : 27 %  
Accuracy of sea : 69 %  
Accuracy of seal : 19 %  
Accuracy of shark : 39 %  
Accuracy of shrew : 14 %  
Accuracy of skunk : 45 %  
Accuracy of skyscraper : 60 %  
Accuracy of snail : 34 %  
Accuracy of snake : 12 %  
Accuracy of spider : 25 %  
Accuracy of squirrel : 15 %  
Accuracy of streetcar : 25 %  
Accuracy of sunflower : 50 %  
Accuracy of sweet\_pepper : 15 %  
Accuracy of table : 23 %  
Accuracy of tank : 46 %  
Accuracy of telephone : 47 %  
Accuracy of television : 49 %  
Accuracy of tiger : 22 %  
Accuracy of tractor : 37 %  
Accuracy of train : 23 %

Accuracy of trout : 33 %  
Accuracy of tulip : 22 %  
Accuracy of turtle : 15 %  
Accuracy of wardrobe : 59 %  
Accuracy of whale : 29 %  
Accuracy of willow\_tree : 13 %  
Accuracy of wolf : 21 %  
Accuracy of woman : 12 %  
Accuracy of worm : 17 %

[12]: *# One of the changes I made was that I added another layer in the network  
# that takes the output from the second convolutional layer, applies ReLU  
↪activation,  
# and passes it through a 2x2 AvgPool layer to capture more relationships.  
# I also reduced the learning rate of the optimizer to half (0.0005) which  
# which allowed the optimizer to take smaller steps towards the minimum of  
# the loss function which might allow for a better chance of finding the global  
↪min of loss.*