

final_14

March 26, 2023

```
[1]: #!pip install torch torchvision
      %matplotlib inline
      import matplotlib.pyplot as plt
      import numpy as np
      import torch
      import torchvision
      import torchvision.transforms as transforms
      import torch.nn as nn
      import torch.nn.functional as F
      import torch.optim as optim
```

Prepare for Dataset

```
[2]: transform_train = transforms.Compose([
      #transforms.RandomCrop(32, padding=4),
      transforms.RandomHorizontalFlip(),
      transforms.RandomRotation(20),
      transforms.ToTensor(),
      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
  ])

  transform_test = transforms.Compose([
      transforms.ToTensor(),
      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
  ])

  trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
      download=True,
      ↪transform=transform_train)
  trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
      shuffle=True, num_workers=2)

  testset = torchvision.datasets.CIFAR100(root='./data', train=False,
      download=True, transform=transform_test)
  testloader = torch.utils.data.DataLoader(testset, batch_size=4,
      shuffle=False, num_workers=2)

  classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
```

```

'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
'castle', 'caterpillar', 'cattle', 'chair',
'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
'elephant', 'flatfish', 'forest', 'fox', 'girl',
'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
'wolf', 'woman', 'worm')

#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

Files already downloaded and verified


Files already downloaded and verified

```

[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))

```



willow_tree skyscraper mouse tulip

Choose a Device

```
[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
padding=1)
        self.bn1 = nn.BatchNorm2d(64)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(128)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(256)
        self.pool3 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu3 = nn.SELU()
        self.fc1 = nn.Linear(256 * 4 * 4, 1024)
        self.bn4 = nn.BatchNorm1d(1024)
        self.relu4 = nn.SELU()
        self.fc2 = nn.Linear(1024, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.conv1(x)))
        x = self.pool2(self.relu2(self.conv2(x)))
```

```

        x = self.pool3(self.relu3(self.conv3(x)))
        x = x.view(-1, 256 * 4 * 4)
        x = self.relu4(self.fc1(x))
        x = self.fc2(x)
        return x

net = Net()      # Create the network instance.
net.to(device)  # Move the network parameters to the specified device.

```

```

[5]: Net(
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu1): ReLU()
  (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool2): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu2): ReLU()
  (conv3): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool3): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu3): SELU()
  (fc1): Linear(in_features=4096, out_features=1024, bias=True)
  (bn4): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu4): SELU()
  (fc2): Linear(in_features=1024, out_features=100, bias=True)
)

```

Optimizer and Loss Function

```

[6]: # We use cross-entropy as loss function.
loss_func = nn.CrossEntropyLoss()
# We use stochastic gradient descent (SGD) as optimizer.
#opt = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)
opt = optim.Adam(net.parameters(), lr=0.0001)

```

Training Procedure

```

[7]: import sys
from tqdm.notebook import tqdm

```

```

avg_losses = []    # Avg. losses.
epochs = 8         # Total epochs.
print_freq = 500   # Print frequency.

for epoch in range(epochs): # Loop over the dataset multiple times.
    running_loss = 0.0      # Initialize running loss.
    for i, data in enumerate(tqdm(trainloader), 0):
        # Get the inputs.
        inputs, labels = data

        # Move the inputs to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients.
        opt.zero_grad()

        # Forward step.
        outputs = net(inputs)
        loss = loss_func(outputs, labels)

        # Backward step.
        loss.backward()

        # Optimization step (update the parameters).
        opt.step()

        # Print statistics.
        running_loss += loss.item()
        if i % print_freq == print_freq - 1: # Print every several mini-batches.
            avg_loss = running_loss / print_freq
            print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.
                ↪format(epoch, i, avg_loss), flush=True)
            sys.stdout.flush()
            avg_losses.append(avg_loss)
            running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 0, i:   499] avg mini-batch loss: 4.406
[epoch: 0, i:   999] avg mini-batch loss: 4.078
[epoch: 0, i:  1499] avg mini-batch loss: 3.967
[epoch: 0, i:  1999] avg mini-batch loss: 3.824
[epoch: 0, i:  2499] avg mini-batch loss: 3.762
[epoch: 0, i:  2999] avg mini-batch loss: 3.662
[epoch: 0, i:  3499] avg mini-batch loss: 3.578
[epoch: 0, i:  3999] avg mini-batch loss: 3.555

```

```

[epoch: 0, i: 4499] avg mini-batch loss: 3.507
[epoch: 0, i: 4999] avg mini-batch loss: 3.407
[epoch: 0, i: 5499] avg mini-batch loss: 3.411
[epoch: 0, i: 5999] avg mini-batch loss: 3.413
[epoch: 0, i: 6499] avg mini-batch loss: 3.284
[epoch: 0, i: 6999] avg mini-batch loss: 3.340
[epoch: 0, i: 7499] avg mini-batch loss: 3.258
[epoch: 0, i: 7999] avg mini-batch loss: 3.243
[epoch: 0, i: 8499] avg mini-batch loss: 3.183
[epoch: 0, i: 8999] avg mini-batch loss: 3.158
[epoch: 0, i: 9499] avg mini-batch loss: 3.160
[epoch: 0, i: 9999] avg mini-batch loss: 3.179
[epoch: 0, i: 10499] avg mini-batch loss: 3.060
[epoch: 0, i: 10999] avg mini-batch loss: 3.025
[epoch: 0, i: 11499] avg mini-batch loss: 3.048
[epoch: 0, i: 11999] avg mini-batch loss: 3.062
[epoch: 0, i: 12499] avg mini-batch loss: 3.007

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 1, i: 499] avg mini-batch loss: 2.921
[epoch: 1, i: 999] avg mini-batch loss: 2.908
[epoch: 1, i: 1499] avg mini-batch loss: 2.927
[epoch: 1, i: 1999] avg mini-batch loss: 2.930
[epoch: 1, i: 2499] avg mini-batch loss: 2.906
[epoch: 1, i: 2999] avg mini-batch loss: 2.876
[epoch: 1, i: 3499] avg mini-batch loss: 2.875
[epoch: 1, i: 3999] avg mini-batch loss: 2.906
[epoch: 1, i: 4499] avg mini-batch loss: 2.838
[epoch: 1, i: 4999] avg mini-batch loss: 2.838
[epoch: 1, i: 5499] avg mini-batch loss: 2.786
[epoch: 1, i: 5999] avg mini-batch loss: 2.861
[epoch: 1, i: 6499] avg mini-batch loss: 2.856
[epoch: 1, i: 6999] avg mini-batch loss: 2.839
[epoch: 1, i: 7499] avg mini-batch loss: 2.848
[epoch: 1, i: 7999] avg mini-batch loss: 2.793
[epoch: 1, i: 8499] avg mini-batch loss: 2.814
[epoch: 1, i: 8999] avg mini-batch loss: 2.721
[epoch: 1, i: 9499] avg mini-batch loss: 2.826
[epoch: 1, i: 9999] avg mini-batch loss: 2.806
[epoch: 1, i: 10499] avg mini-batch loss: 2.739
[epoch: 1, i: 10999] avg mini-batch loss: 2.807
[epoch: 1, i: 11499] avg mini-batch loss: 2.717
[epoch: 1, i: 11999] avg mini-batch loss: 2.707
[epoch: 1, i: 12499] avg mini-batch loss: 2.809

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 2, i: 499] avg mini-batch loss: 2.663
[epoch: 2, i: 999] avg mini-batch loss: 2.648

```

```

[epoch: 2, i: 1499] avg mini-batch loss: 2.638
[epoch: 2, i: 1999] avg mini-batch loss: 2.603
[epoch: 2, i: 2499] avg mini-batch loss: 2.597
[epoch: 2, i: 2999] avg mini-batch loss: 2.612
[epoch: 2, i: 3499] avg mini-batch loss: 2.605
[epoch: 2, i: 3999] avg mini-batch loss: 2.602
[epoch: 2, i: 4499] avg mini-batch loss: 2.574
[epoch: 2, i: 4999] avg mini-batch loss: 2.612
[epoch: 2, i: 5499] avg mini-batch loss: 2.666
[epoch: 2, i: 5999] avg mini-batch loss: 2.620
[epoch: 2, i: 6499] avg mini-batch loss: 2.595
[epoch: 2, i: 6999] avg mini-batch loss: 2.623
[epoch: 2, i: 7499] avg mini-batch loss: 2.596
[epoch: 2, i: 7999] avg mini-batch loss: 2.643
[epoch: 2, i: 8499] avg mini-batch loss: 2.566
[epoch: 2, i: 8999] avg mini-batch loss: 2.579
[epoch: 2, i: 9499] avg mini-batch loss: 2.534
[epoch: 2, i: 9999] avg mini-batch loss: 2.626
[epoch: 2, i: 10499] avg mini-batch loss: 2.572
[epoch: 2, i: 10999] avg mini-batch loss: 2.610
[epoch: 2, i: 11499] avg mini-batch loss: 2.570
[epoch: 2, i: 11999] avg mini-batch loss: 2.563
[epoch: 2, i: 12499] avg mini-batch loss: 2.611

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 3, i: 499] avg mini-batch loss: 2.499
[epoch: 3, i: 999] avg mini-batch loss: 2.478
[epoch: 3, i: 1499] avg mini-batch loss: 2.476
[epoch: 3, i: 1999] avg mini-batch loss: 2.451
[epoch: 3, i: 2499] avg mini-batch loss: 2.477
[epoch: 3, i: 2999] avg mini-batch loss: 2.463
[epoch: 3, i: 3499] avg mini-batch loss: 2.539
[epoch: 3, i: 3999] avg mini-batch loss: 2.401
[epoch: 3, i: 4499] avg mini-batch loss: 2.482
[epoch: 3, i: 4999] avg mini-batch loss: 2.451
[epoch: 3, i: 5499] avg mini-batch loss: 2.383
[epoch: 3, i: 5999] avg mini-batch loss: 2.401
[epoch: 3, i: 6499] avg mini-batch loss: 2.468
[epoch: 3, i: 6999] avg mini-batch loss: 2.368
[epoch: 3, i: 7499] avg mini-batch loss: 2.430
[epoch: 3, i: 7999] avg mini-batch loss: 2.415
[epoch: 3, i: 8499] avg mini-batch loss: 2.466
[epoch: 3, i: 8999] avg mini-batch loss: 2.478
[epoch: 3, i: 9499] avg mini-batch loss: 2.451
[epoch: 3, i: 9999] avg mini-batch loss: 2.455
[epoch: 3, i: 10499] avg mini-batch loss: 2.392
[epoch: 3, i: 10999] avg mini-batch loss: 2.426
[epoch: 3, i: 11499] avg mini-batch loss: 2.505

```

[epoch: 3, i: 11999] avg mini-batch loss: 2.401
[epoch: 3, i: 12499] avg mini-batch loss: 2.341

0%| | 0/12500 [00:00<?, ?it/s]

[epoch: 4, i: 499] avg mini-batch loss: 2.213
[epoch: 4, i: 999] avg mini-batch loss: 2.363
[epoch: 4, i: 1499] avg mini-batch loss: 2.302
[epoch: 4, i: 1999] avg mini-batch loss: 2.337
[epoch: 4, i: 2499] avg mini-batch loss: 2.372
[epoch: 4, i: 2999] avg mini-batch loss: 2.391
[epoch: 4, i: 3499] avg mini-batch loss: 2.372
[epoch: 4, i: 3999] avg mini-batch loss: 2.372
[epoch: 4, i: 4499] avg mini-batch loss: 2.362
[epoch: 4, i: 4999] avg mini-batch loss: 2.274
[epoch: 4, i: 5499] avg mini-batch loss: 2.331
[epoch: 4, i: 5999] avg mini-batch loss: 2.304
[epoch: 4, i: 6499] avg mini-batch loss: 2.406
[epoch: 4, i: 6999] avg mini-batch loss: 2.332
[epoch: 4, i: 7499] avg mini-batch loss: 2.281
[epoch: 4, i: 7999] avg mini-batch loss: 2.405
[epoch: 4, i: 8499] avg mini-batch loss: 2.292
[epoch: 4, i: 8999] avg mini-batch loss: 2.306
[epoch: 4, i: 9499] avg mini-batch loss: 2.325
[epoch: 4, i: 9999] avg mini-batch loss: 2.238
[epoch: 4, i: 10499] avg mini-batch loss: 2.283
[epoch: 4, i: 10999] avg mini-batch loss: 2.316
[epoch: 4, i: 11499] avg mini-batch loss: 2.287
[epoch: 4, i: 11999] avg mini-batch loss: 2.329
[epoch: 4, i: 12499] avg mini-batch loss: 2.318

0%| | 0/12500 [00:00<?, ?it/s]

[epoch: 5, i: 499] avg mini-batch loss: 2.210
[epoch: 5, i: 999] avg mini-batch loss: 2.197
[epoch: 5, i: 1499] avg mini-batch loss: 2.162
[epoch: 5, i: 1999] avg mini-batch loss: 2.244
[epoch: 5, i: 2499] avg mini-batch loss: 2.256
[epoch: 5, i: 2999] avg mini-batch loss: 2.237
[epoch: 5, i: 3499] avg mini-batch loss: 2.176
[epoch: 5, i: 3999] avg mini-batch loss: 2.211
[epoch: 5, i: 4499] avg mini-batch loss: 2.139
[epoch: 5, i: 4999] avg mini-batch loss: 2.281
[epoch: 5, i: 5499] avg mini-batch loss: 2.183
[epoch: 5, i: 5999] avg mini-batch loss: 2.276
[epoch: 5, i: 6499] avg mini-batch loss: 2.143
[epoch: 5, i: 6999] avg mini-batch loss: 2.215
[epoch: 5, i: 7499] avg mini-batch loss: 2.165
[epoch: 5, i: 7999] avg mini-batch loss: 2.225
[epoch: 5, i: 8499] avg mini-batch loss: 2.316


```

[epoch: 5, i: 8999] avg mini-batch loss: 2.138
[epoch: 5, i: 9499] avg mini-batch loss: 2.220
[epoch: 5, i: 9999] avg mini-batch loss: 2.246
[epoch: 5, i: 10499] avg mini-batch loss: 2.224
[epoch: 5, i: 10999] avg mini-batch loss: 2.254
[epoch: 5, i: 11499] avg mini-batch loss: 2.195
[epoch: 5, i: 11999] avg mini-batch loss: 2.212
[epoch: 5, i: 12499] avg mini-batch loss: 2.218

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 6, i: 499] avg mini-batch loss: 2.055
[epoch: 6, i: 999] avg mini-batch loss: 2.136
[epoch: 6, i: 1499] avg mini-batch loss: 2.063
[epoch: 6, i: 1999] avg mini-batch loss: 2.153
[epoch: 6, i: 2499] avg mini-batch loss: 2.105
[epoch: 6, i: 2999] avg mini-batch loss: 2.074
[epoch: 6, i: 3499] avg mini-batch loss: 2.157
[epoch: 6, i: 3999] avg mini-batch loss: 2.177
[epoch: 6, i: 4499] avg mini-batch loss: 2.100
[epoch: 6, i: 4999] avg mini-batch loss: 2.127
[epoch: 6, i: 5499] avg mini-batch loss: 2.116
[epoch: 6, i: 5999] avg mini-batch loss: 2.164
[epoch: 6, i: 6499] avg mini-batch loss: 2.115
[epoch: 6, i: 6999] avg mini-batch loss: 2.081
[epoch: 6, i: 7499] avg mini-batch loss: 2.180
[epoch: 6, i: 7999] avg mini-batch loss: 2.152
[epoch: 6, i: 8499] avg mini-batch loss: 2.041
[epoch: 6, i: 8999] avg mini-batch loss: 2.101
[epoch: 6, i: 9499] avg mini-batch loss: 2.129
[epoch: 6, i: 9999] avg mini-batch loss: 2.155
[epoch: 6, i: 10499] avg mini-batch loss: 2.067
[epoch: 6, i: 10999] avg mini-batch loss: 2.228
[epoch: 6, i: 11499] avg mini-batch loss: 2.108
[epoch: 6, i: 11999] avg mini-batch loss: 2.167
[epoch: 6, i: 12499] avg mini-batch loss: 2.126

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```


[epoch: 7, i: 499] avg mini-batch loss: 2.062
[epoch: 7, i: 999] avg mini-batch loss: 2.056
[epoch: 7, i: 1499] avg mini-batch loss: 2.006
[epoch: 7, i: 1999] avg mini-batch loss: 2.060
[epoch: 7, i: 2499] avg mini-batch loss: 2.005
[epoch: 7, i: 2999] avg mini-batch loss: 2.052
[epoch: 7, i: 3499] avg mini-batch loss: 2.068
[epoch: 7, i: 3999] avg mini-batch loss: 2.066
[epoch: 7, i: 4499] avg mini-batch loss: 1.968
[epoch: 7, i: 4999] avg mini-batch loss: 2.023
[epoch: 7, i: 5499] avg mini-batch loss: 2.000

```

```
[epoch: 7, i: 5999] avg mini-batch loss: 2.046
[epoch: 7, i: 6499] avg mini-batch loss: 2.024
[epoch: 7, i: 6999] avg mini-batch loss: 2.043
[epoch: 7, i: 7499] avg mini-batch loss: 2.064
[epoch: 7, i: 7999] avg mini-batch loss: 2.012
[epoch: 7, i: 8499] avg mini-batch loss: 2.025
[epoch: 7, i: 8999] avg mini-batch loss: 2.039
[epoch: 7, i: 9499] avg mini-batch loss: 2.103
[epoch: 7, i: 9999] avg mini-batch loss: 2.022
[epoch: 7, i: 10499] avg mini-batch loss: 2.025
[epoch: 7, i: 10999] avg mini-batch loss: 2.075
[epoch: 7, i: 11499] avg mini-batch loss: 2.040
[epoch: 7, i: 11999] avg mini-batch loss: 2.040
[epoch: 7, i: 12499] avg mini-batch loss: 2.047
Finished Training.
```

Training Loss Curve

```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                              for j in range(4)))
```



GroundTruth: mountain forest seal mushroom
Predicted: road squirrel caterpillar mushroom

```
[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 45 %

```
[11]: # Get test accuracy for each class.
class_correct = [0] * len(classes)
class_total = [0] * len(classes)
with torch.no_grad():
```

```

for data in testloader:
    images, labels = data
    images, labels = images.to(device), labels.to(device)
    outputs = net(images)
    _, predicted = torch.max(outputs, 1)
    c = (predicted == labels).squeeze()
    for i in range(len(labels)):
        label = labels[i]
        class_correct[label] += c[i].item()
        class_total[label] += 1

for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

```

Accuracy of apple : 67 %
Accuracy of aquarium_fish : 69 %
Accuracy of  baby : 37 %
Accuracy of  bear : 18 %
Accuracy of beaver : 27 %
Accuracy of  bed : 42 %
Accuracy of  bee : 54 %
Accuracy of beetle : 42 %
Accuracy of bicycle : 50 %
Accuracy of bottle : 61 %
Accuracy of  bowl : 20 %
Accuracy of  boy : 24 %
Accuracy of bridge : 42 %
Accuracy of  bus : 40 %
Accuracy of butterfly : 43 %
Accuracy of camel : 21 %
Accuracy of  can : 45 %
Accuracy of castle : 55 %
Accuracy of caterpillar : 52 %
Accuracy of cattle : 40 %
Accuracy of chair : 63 %
Accuracy of chimpanzee : 84 %
Accuracy of clock : 41 %
Accuracy of cloud : 74 %
Accuracy of cockroach : 69 %
Accuracy of couch : 39 %
Accuracy of  crab : 41 %
Accuracy of crocodile : 31 %
Accuracy of  cup : 67 %
Accuracy of dinosaur : 38 %
Accuracy of dolphin : 38 %
Accuracy of elephant : 36 %

```

Accuracy of flatfish : 46 %
Accuracy of forest : 38 %
Accuracy of fox : 41 %
Accuracy of girl : 28 %
Accuracy of hamster : 40 %
Accuracy of house : 63 %
Accuracy of kangaroo : 19 %
Accuracy of keyboard : 56 %
Accuracy of lamp : 35 %
Accuracy of lawn_mower : 70 %
Accuracy of leopard : 25 %
Accuracy of lion : 48 %
Accuracy of lizard : 19 %
Accuracy of lobster : 34 %
Accuracy of man : 21 %
Accuracy of maple_tree : 61 %
Accuracy of motorcycle : 81 %
Accuracy of mountain : 61 %
Accuracy of mouse : 36 %
Accuracy of mushroom : 42 %
Accuracy of oak_tree : 54 %
Accuracy of orange : 60 %
Accuracy of orchid : 70 %
Accuracy of otter : 6 %
Accuracy of palm_tree : 63 %
Accuracy of pear : 49 %
Accuracy of pickup_truck : 60 %
Accuracy of pine_tree : 42 %
Accuracy of plain : 73 %
Accuracy of plate : 48 %
Accuracy of poppy : 31 %
Accuracy of porcupine : 39 %
Accuracy of possum : 29 %
Accuracy of rabbit : 19 %
Accuracy of raccoon : 29 %
Accuracy of ray : 36 %
Accuracy of road : 80 %
Accuracy of rocket : 62 %
Accuracy of rose : 61 %
Accuracy of sea : 62 %
Accuracy of seal : 7 %
Accuracy of shark : 46 %
Accuracy of shrew : 34 %
Accuracy of skunk : 62 %
Accuracy of skyscraper : 75 %
Accuracy of snail : 20 %
Accuracy of snake : 42 %
Accuracy of spider : 41 %

Accuracy of squirrel : 23 %
Accuracy of streetcar : 61 %
Accuracy of sunflower : 76 %
Accuracy of sweet_pepper : 51 %
Accuracy of table : 48 %
Accuracy of tank : 48 %
Accuracy of telephone : 48 %
Accuracy of television : 50 %
Accuracy of tiger : 47 %
Accuracy of tractor : 54 %
Accuracy of train : 40 %
Accuracy of trout : 54 %
Accuracy of tulip : 22 %
Accuracy of turtle : 31 %
Accuracy of wardrobe : 75 %
Accuracy of whale : 48 %
Accuracy of willow_tree : 35 %
Accuracy of wolf : 51 %
Accuracy of woman : 24 %
Accuracy of worm : 41 %

```
[12]: # One of the changes I made was that I added another layer in the network
      # that takes the output from the second convolutional layer, applies ReLU
      ↪activation,
      # and passes it through a 2x2 AvgPool layer to capture more relationships.
      # I also reduced the learning rate of the optimizer to half (0.0005) which
      # which allowed the optimizer to take smaller steps towards the minimum of
      # the loss function which might allow for a better chance of finding the global
      ↪min of loss.
```