

final_21

March 26, 2023

```
[1]: #!pip install torch torchvision
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

Prepare for Dataset

```
[2]: transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
           'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
           'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
           'castle', 'caterpillar', 'cattle', 'chair',
           'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
```

```
'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
'elephant', 'flatfish', 'forest', 'fox', 'girl',
'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
'wolf', 'woman', 'worm')
```


```
#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



girl rabbit skunk dinosaur

Choose a Device

```
[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
padding=1)
        self.bn1 = nn.BatchNorm2d(64)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(128)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(256)
        self.pool3 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu3 = nn.ReLU()
        self.conv4 = nn.Conv2d(256, 512, 3, padding=1)
        self.bn4 = nn.BatchNorm2d(512)
        self.pool4 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu4 = nn.ReLU()
        self.fc1 = nn.Linear(512 * 2 * 2, 1024)
        self.bn5 = nn.BatchNorm1d(1024)
        self.relu5 = nn.ReLU()
        self.fc2 = nn.Linear(1024, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.conv1(x)))
        x = self.pool2(self.relu2(self.conv2(x)))
        x = self.pool3(self.relu3(self.conv3(x)))
        x = self.pool4(self.relu4(self.bn4(self.conv4(x))))
        x = x.view(-1, 512 * 2 * 2)
        x = self.relu5(self.bn5(self.fc1(x)))
        x = self.fc2(x)
        return x
```

```
net = Net()      # Create the network instance.
net.to(device)  # Move the network parameters to the specified device.
```

```
[5]: Net(
    (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool1): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (relu1): ReLU()
    (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool2): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (relu2): ReLU()
    (conv3): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool3): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (relu3): ReLU()
    (conv4): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (pool4): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (relu4): ReLU()
    (fc1): Linear(in_features=2048, out_features=1024, bias=True)
    (bn5): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu5): ReLU()
    (fc2): Linear(in_features=1024, out_features=100, bias=True)
)
```

Optimizer and Loss Function

```
[6]: # We use cross-entropy as loss function.
loss_func = nn.CrossEntropyLoss()
# We use stochastic gradient descent (SGD) as optimizer.
#opt = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)
opt = optim.Adam(net.parameters(), lr=0.0001)
```

Training Procedure

```
[7]: import sys
from tqdm.notebook import tqdm

avg_losses = []    # Avg. losses.
epochs = 9         # Total epochs.
```

```

print_freq = 500 # Print frequency.

for epoch in range(epochs): # Loop over the dataset multiple times.
    running_loss = 0.0 # Initialize running loss.
    for i, data in enumerate(tqdm(trainloader), 0):
        # Get the inputs.
        inputs, labels = data

        # Move the inputs to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients.
        opt.zero_grad()

        # Forward step.
        outputs = net(inputs)
        loss = loss_func(outputs, labels)

        # Backward step.
        loss.backward()

        # Optimization step (update the parameters).
        opt.step()

        # Print statistics.
        running_loss += loss.item()
        if i % print_freq == print_freq - 1: # Print every several mini-batches.
            avg_loss = running_loss / print_freq
            print('[epoch: {}, i: {:5d}] avg mini-batch loss: {:.3f}'.
                ↪format(epoch, i, avg_loss), flush=True)
            sys.stdout.flush()
            avg_losses.append(avg_loss)
            running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 0, i: 499] avg mini-batch loss: 4.509
[epoch: 0, i: 999] avg mini-batch loss: 4.274
[epoch: 0, i: 1499] avg mini-batch loss: 4.172
[epoch: 0, i: 1999] avg mini-batch loss: 4.147
[epoch: 0, i: 2499] avg mini-batch loss: 4.068
[epoch: 0, i: 2999] avg mini-batch loss: 4.022
[epoch: 0, i: 3499] avg mini-batch loss: 3.974
[epoch: 0, i: 3999] avg mini-batch loss: 3.863
[epoch: 0, i: 4499] avg mini-batch loss: 3.837
[epoch: 0, i: 4999] avg mini-batch loss: 3.793

```

```

[epoch: 0, i: 5499] avg mini-batch loss: 3.768
[epoch: 0, i: 5999] avg mini-batch loss: 3.713
[epoch: 0, i: 6499] avg mini-batch loss: 3.647
[epoch: 0, i: 6999] avg mini-batch loss: 3.588
[epoch: 0, i: 7499] avg mini-batch loss: 3.548
[epoch: 0, i: 7999] avg mini-batch loss: 3.607
[epoch: 0, i: 8499] avg mini-batch loss: 3.550
[epoch: 0, i: 8999] avg mini-batch loss: 3.480
[epoch: 0, i: 9499] avg mini-batch loss: 3.465
[epoch: 0, i: 9999] avg mini-batch loss: 3.458
[epoch: 0, i: 10499] avg mini-batch loss: 3.412
[epoch: 0, i: 10999] avg mini-batch loss: 3.320
[epoch: 0, i: 11499] avg mini-batch loss: 3.361
[epoch: 0, i: 11999] avg mini-batch loss: 3.353
[epoch: 0, i: 12499] avg mini-batch loss: 3.240

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 1, i: 499] avg mini-batch loss: 3.191
[epoch: 1, i: 999] avg mini-batch loss: 3.161
[epoch: 1, i: 1499] avg mini-batch loss: 3.121
[epoch: 1, i: 1999] avg mini-batch loss: 3.129
[epoch: 1, i: 2499] avg mini-batch loss: 3.112
[epoch: 1, i: 2999] avg mini-batch loss: 3.083
[epoch: 1, i: 3499] avg mini-batch loss: 3.024
[epoch: 1, i: 3999] avg mini-batch loss: 3.083
[epoch: 1, i: 4499] avg mini-batch loss: 3.034
[epoch: 1, i: 4999] avg mini-batch loss: 2.926
[epoch: 1, i: 5499] avg mini-batch loss: 3.057
[epoch: 1, i: 5999] avg mini-batch loss: 2.997
[epoch: 1, i: 6499] avg mini-batch loss: 2.936
[epoch: 1, i: 6999] avg mini-batch loss: 2.912
[epoch: 1, i: 7499] avg mini-batch loss: 2.998
[epoch: 1, i: 7999] avg mini-batch loss: 2.902
[epoch: 1, i: 8499] avg mini-batch loss: 2.911
[epoch: 1, i: 8999] avg mini-batch loss: 2.943
[epoch: 1, i: 9499] avg mini-batch loss: 2.886
[epoch: 1, i: 9999] avg mini-batch loss: 2.876
[epoch: 1, i: 10499] avg mini-batch loss: 2.850
[epoch: 1, i: 10999] avg mini-batch loss: 2.767
[epoch: 1, i: 11499] avg mini-batch loss: 2.771
[epoch: 1, i: 11999] avg mini-batch loss: 2.841
[epoch: 1, i: 12499] avg mini-batch loss: 2.799

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 2, i: 499] avg mini-batch loss: 2.613
[epoch: 2, i: 999] avg mini-batch loss: 2.662
[epoch: 2, i: 1499] avg mini-batch loss: 2.561
[epoch: 2, i: 1999] avg mini-batch loss: 2.655

```

```

[epoch: 2, i: 2499] avg mini-batch loss: 2.665
[epoch: 2, i: 2999] avg mini-batch loss: 2.705
[epoch: 2, i: 3499] avg mini-batch loss: 2.624
[epoch: 2, i: 3999] avg mini-batch loss: 2.584
[epoch: 2, i: 4499] avg mini-batch loss: 2.642
[epoch: 2, i: 4999] avg mini-batch loss: 2.605
[epoch: 2, i: 5499] avg mini-batch loss: 2.631
[epoch: 2, i: 5999] avg mini-batch loss: 2.633
[epoch: 2, i: 6499] avg mini-batch loss: 2.573
[epoch: 2, i: 6999] avg mini-batch loss: 2.557
[epoch: 2, i: 7499] avg mini-batch loss: 2.559
[epoch: 2, i: 7999] avg mini-batch loss: 2.603
[epoch: 2, i: 8499] avg mini-batch loss: 2.559
[epoch: 2, i: 8999] avg mini-batch loss: 2.553
[epoch: 2, i: 9499] avg mini-batch loss: 2.573
[epoch: 2, i: 9999] avg mini-batch loss: 2.450
[epoch: 2, i: 10499] avg mini-batch loss: 2.609
[epoch: 2, i: 10999] avg mini-batch loss: 2.443
[epoch: 2, i: 11499] avg mini-batch loss: 2.515
[epoch: 2, i: 11999] avg mini-batch loss: 2.465
[epoch: 2, i: 12499] avg mini-batch loss: 2.428

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 3, i: 499] avg mini-batch loss: 2.338
[epoch: 3, i: 999] avg mini-batch loss: 2.311
[epoch: 3, i: 1499] avg mini-batch loss: 2.368
[epoch: 3, i: 1999] avg mini-batch loss: 2.360
[epoch: 3, i: 2499] avg mini-batch loss: 2.356
[epoch: 3, i: 2999] avg mini-batch loss: 2.315
[epoch: 3, i: 3499] avg mini-batch loss: 2.344
[epoch: 3, i: 3999] avg mini-batch loss: 2.339
[epoch: 3, i: 4499] avg mini-batch loss: 2.422
[epoch: 3, i: 4999] avg mini-batch loss: 2.260
[epoch: 3, i: 5499] avg mini-batch loss: 2.330
[epoch: 3, i: 5999] avg mini-batch loss: 2.297
[epoch: 3, i: 6499] avg mini-batch loss: 2.317
[epoch: 3, i: 6999] avg mini-batch loss: 2.327
[epoch: 3, i: 7499] avg mini-batch loss: 2.293
[epoch: 3, i: 7999] avg mini-batch loss: 2.310
[epoch: 3, i: 8499] avg mini-batch loss: 2.265
[epoch: 3, i: 8999] avg mini-batch loss: 2.265
[epoch: 3, i: 9499] avg mini-batch loss: 2.336
[epoch: 3, i: 9999] avg mini-batch loss: 2.293
[epoch: 3, i: 10499] avg mini-batch loss: 2.243
[epoch: 3, i: 10999] avg mini-batch loss: 2.274
[epoch: 3, i: 11499] avg mini-batch loss: 2.271
[epoch: 3, i: 11999] avg mini-batch loss: 2.354
[epoch: 3, i: 12499] avg mini-batch loss: 2.279

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 4, i: 499] avg mini-batch loss: 2.020
[epoch: 4, i: 999] avg mini-batch loss: 2.074
[epoch: 4, i: 1499] avg mini-batch loss: 2.084
[epoch: 4, i: 1999] avg mini-batch loss: 2.038
[epoch: 4, i: 2499] avg mini-batch loss: 2.084
[epoch: 4, i: 2999] avg mini-batch loss: 2.168
[epoch: 4, i: 3499] avg mini-batch loss: 2.066
[epoch: 4, i: 3999] avg mini-batch loss: 2.096
[epoch: 4, i: 4499] avg mini-batch loss: 2.109
[epoch: 4, i: 4999] avg mini-batch loss: 2.056
[epoch: 4, i: 5499] avg mini-batch loss: 2.060
[epoch: 4, i: 5999] avg mini-batch loss: 2.078
[epoch: 4, i: 6499] avg mini-batch loss: 2.068
[epoch: 4, i: 6999] avg mini-batch loss: 2.168
[epoch: 4, i: 7499] avg mini-batch loss: 2.109
[epoch: 4, i: 7999] avg mini-batch loss: 2.062
[epoch: 4, i: 8499] avg mini-batch loss: 2.101
[epoch: 4, i: 8999] avg mini-batch loss: 2.070
[epoch: 4, i: 9499] avg mini-batch loss: 2.126
[epoch: 4, i: 9999] avg mini-batch loss: 2.092
[epoch: 4, i: 10499] avg mini-batch loss: 2.021
[epoch: 4, i: 10999] avg mini-batch loss: 2.112
[epoch: 4, i: 11499] avg mini-batch loss: 2.100
[epoch: 4, i: 11999] avg mini-batch loss: 2.077
[epoch: 4, i: 12499] avg mini-batch loss: 2.031

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 5, i: 499] avg mini-batch loss: 1.883
[epoch: 5, i: 999] avg mini-batch loss: 1.947
[epoch: 5, i: 1499] avg mini-batch loss: 1.822
[epoch: 5, i: 1999] avg mini-batch loss: 1.862
[epoch: 5, i: 2499] avg mini-batch loss: 1.789
[epoch: 5, i: 2999] avg mini-batch loss: 1.874
[epoch: 5, i: 3499] avg mini-batch loss: 1.908
[epoch: 5, i: 3999] avg mini-batch loss: 1.827
[epoch: 5, i: 4499] avg mini-batch loss: 1.922
[epoch: 5, i: 4999] avg mini-batch loss: 1.879
[epoch: 5, i: 5499] avg mini-batch loss: 1.873
[epoch: 5, i: 5999] avg mini-batch loss: 1.906
[epoch: 5, i: 6499] avg mini-batch loss: 1.852
[epoch: 5, i: 6999] avg mini-batch loss: 1.985
[epoch: 5, i: 7499] avg mini-batch loss: 1.891
[epoch: 5, i: 7999] avg mini-batch loss: 1.856
[epoch: 5, i: 8499] avg mini-batch loss: 1.888
[epoch: 5, i: 8999] avg mini-batch loss: 1.891
[epoch: 5, i: 9499] avg mini-batch loss: 1.946

```



```

[epoch: 5, i: 9999] avg mini-batch loss: 1.866
[epoch: 5, i: 10499] avg mini-batch loss: 1.849
[epoch: 5, i: 10999] avg mini-batch loss: 1.926
[epoch: 5, i: 11499] avg mini-batch loss: 1.918
[epoch: 5, i: 11999] avg mini-batch loss: 1.920
[epoch: 5, i: 12499] avg mini-batch loss: 1.949

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 6, i: 499] avg mini-batch loss: 1.657
[epoch: 6, i: 999] avg mini-batch loss: 1.725
[epoch: 6, i: 1499] avg mini-batch loss: 1.620
[epoch: 6, i: 1999] avg mini-batch loss: 1.738
[epoch: 6, i: 2499] avg mini-batch loss: 1.704
[epoch: 6, i: 2999] avg mini-batch loss: 1.643
[epoch: 6, i: 3499] avg mini-batch loss: 1.688
[epoch: 6, i: 3999] avg mini-batch loss: 1.708
[epoch: 6, i: 4499] avg mini-batch loss: 1.760
[epoch: 6, i: 4999] avg mini-batch loss: 1.691
[epoch: 6, i: 5499] avg mini-batch loss: 1.715
[epoch: 6, i: 5999] avg mini-batch loss: 1.734
[epoch: 6, i: 6499] avg mini-batch loss: 1.704
[epoch: 6, i: 6999] avg mini-batch loss: 1.729
[epoch: 6, i: 7499] avg mini-batch loss: 1.730
[epoch: 6, i: 7999] avg mini-batch loss: 1.684
[epoch: 6, i: 8499] avg mini-batch loss: 1.741
[epoch: 6, i: 8999] avg mini-batch loss: 1.683
[epoch: 6, i: 9499] avg mini-batch loss: 1.733
[epoch: 6, i: 9999] avg mini-batch loss: 1.765
[epoch: 6, i: 10499] avg mini-batch loss: 1.699
[epoch: 6, i: 10999] avg mini-batch loss: 1.702
[epoch: 6, i: 11499] avg mini-batch loss: 1.697
[epoch: 6, i: 11999] avg mini-batch loss: 1.743
[epoch: 6, i: 12499] avg mini-batch loss: 1.730

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 7, i: 499] avg mini-batch loss: 1.498
[epoch: 7, i: 999] avg mini-batch loss: 1.489
[epoch: 7, i: 1499] avg mini-batch loss: 1.537
[epoch: 7, i: 1999] avg mini-batch loss: 1.551
[epoch: 7, i: 2499] avg mini-batch loss: 1.506
[epoch: 7, i: 2999] avg mini-batch loss: 1.455
[epoch: 7, i: 3499] avg mini-batch loss: 1.547
[epoch: 7, i: 3999] avg mini-batch loss: 1.548
[epoch: 7, i: 4499] avg mini-batch loss: 1.514
[epoch: 7, i: 4999] avg mini-batch loss: 1.487
[epoch: 7, i: 5499] avg mini-batch loss: 1.491
[epoch: 7, i: 5999] avg mini-batch loss: 1.541
[epoch: 7, i: 6499] avg mini-batch loss: 1.515

```

```
[epoch: 7, i: 6999] avg mini-batch loss: 1.560
```

IOPub message rate exceeded.

The notebook server will temporarily stop sending output to the client in order to avoid crashing it.

To change this limit, set the config variable

```
`--NotebookApp.iopub_msg_rate_limit`.
```


Current values:

NotebookApp.iopub_msg_rate_limit=1000.0 (msgs/sec)

NotebookApp.rate_limit_window=3.0 (secs)

Training Loss Curve

```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



Evaluate on Test Dataset


```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
```

```

print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                               for j in range(4)))

```



GroundTruth: mountain forest seal mushroom
Predicted: train rabbit dolphin mushroom

```

[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))

```

Accuracy of the network on the 10000 test images: 46 %

```

[11]: # Get test accuracy for each class.
class_correct = [0] * len(classes)
class_total = [0] * len(classes)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)

```

```

_, predicted = torch.max(outputs, 1)
c = (predicted == labels).squeeze()
for i in range(len(labels)):
    label = labels[i]
    class_correct[label] += c[i].item()
    class_total[label] += 1

for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

```

Accuracy of apple : 78 %
Accuracy of aquarium_fish : 53 %
Accuracy of  baby : 21 %
Accuracy of  bear : 23 %
Accuracy of beaver : 20 %
Accuracy of  bed : 34 %
Accuracy of  bee : 70 %
Accuracy of beetle : 48 %
Accuracy of bicycle : 55 %
Accuracy of bottle : 62 %
Accuracy of  bowl : 39 %
Accuracy of  boy : 18 %
Accuracy of bridge : 64 %
Accuracy of  bus : 40 %
Accuracy of butterfly : 31 %
Accuracy of camel : 33 %
Accuracy of  can : 53 %
Accuracy of castle : 60 %
Accuracy of caterpillar : 42 %
Accuracy of cattle : 35 %
Accuracy of chair : 73 %
Accuracy of chimpanzee : 62 %
Accuracy of clock : 45 %
Accuracy of cloud : 77 %
Accuracy of cockroach : 59 %
Accuracy of couch : 42 %
Accuracy of  crab : 42 %
Accuracy of crocodile : 30 %
Accuracy of  cup : 65 %
Accuracy of dinosaur : 40 %
Accuracy of dolphin : 50 %
Accuracy of elephant : 52 %
Accuracy of flatfish : 40 %
Accuracy of forest : 51 %
Accuracy of  fox : 32 %
Accuracy of  girl : 62 %

```

Accuracy of hamster : 39 %
Accuracy of house : 43 %
Accuracy of kangaroo : 18 %
Accuracy of keyboard : 56 %
Accuracy of lamp : 40 %
Accuracy of lawn_mower : 66 %
Accuracy of leopard : 41 %
Accuracy of lion : 49 %
Accuracy of lizard : 15 %
Accuracy of lobster : 25 %
Accuracy of man : 22 %
Accuracy of maple_tree : 76 %
Accuracy of motorcycle : 74 %
Accuracy of mountain : 45 %
Accuracy of mouse : 21 %
Accuracy of mushroom : 44 %
Accuracy of oak_tree : 31 %
Accuracy of orange : 58 %
Accuracy of orchid : 70 %
Accuracy of otter : 6 %
Accuracy of palm_tree : 60 %
Accuracy of pear : 58 %
Accuracy of pickup_truck : 76 %
Accuracy of pine_tree : 43 %
Accuracy of plain : 75 %
Accuracy of plate : 57 %
Accuracy of poppy : 54 %
Accuracy of porcupine : 44 %
Accuracy of possum : 38 %
Accuracy of rabbit : 23 %
Accuracy of raccoon : 38 %
Accuracy of ray : 31 %
Accuracy of road : 81 %
Accuracy of rocket : 59 %
Accuracy of rose : 43 %
Accuracy of sea : 58 %
Accuracy of seal : 16 %
Accuracy of shark : 27 %
Accuracy of shrew : 20 %
Accuracy of skunk : 71 %
Accuracy of skyscraper : 69 %
Accuracy of snail : 31 %
Accuracy of snake : 24 %
Accuracy of spider : 38 %
Accuracy of squirrel : 19 %
Accuracy of streetcar : 43 %
Accuracy of sunflower : 73 %
Accuracy of sweet_pepper : 45 %

Accuracy of table : 49 %
Accuracy of tank : 54 %
Accuracy of telephone : 48 %
Accuracy of television : 52 %
Accuracy of tiger : 42 %
Accuracy of tractor : 67 %
Accuracy of train : 53 %
Accuracy of trout : 60 %
Accuracy of tulip : 31 %
Accuracy of turtle : 38 %
Accuracy of wardrobe : 78 %
Accuracy of whale : 64 %
Accuracy of willow_tree : 24 %
Accuracy of wolf : 45 %
Accuracy of woman : 11 %
Accuracy of worm : 44 %

```
[12]: # One of the changes I made was that I added another layer in the network
      # that takes the output from the second convolutional layer, applies ReLU
      ↪activation,
      # and passes it through a 2x2 AvgPool layer to capture more relationships.
      # I also reduced the learning rate of the optimizer to half (0.0005) which
      # which allowed the optimizer to take smaller steps towards the minimum of
      # the loss function which might allow for a better chance of finding the global
      ↪min of loss.
```