

## final\_23

March 26, 2023

```
[1]: #!pip install torch torchvision
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

### Prepare for Dataset

```
[11]: transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
           'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
           'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
           'castle', 'caterpillar', 'cattle', 'chair',
           'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
```

```
'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
'elephant', 'flatfish', 'forest', 'fox', 'girl',
'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
'wolf', 'woman', 'worm')
```


```
#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



spider rabbit kangaroo chair

### Choose a Device

```
[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

cuda:0

### Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=128, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(128)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(256)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(256, 512, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(512)
        self.pool3 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu3 = nn.ReLU()
        self.conv4 = nn.Conv2d(512, 1024, 3, padding=1)
        self.bn4 = nn.BatchNorm2d(1024)
        self.pool4 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu4 = nn.ReLU()
        self.fc1 = nn.Linear(1024 * 2 * 2, 2048)
        self.bn5 = nn.BatchNorm1d(2048)
        self.relu5 = nn.ReLU()
        self.fc2 = nn.Linear(2048, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.bn1(self.conv1(x))))
        x = self.pool2(self.relu2(self.bn2(self.conv2(x))))
        x = self.pool3(self.relu3(self.bn3(self.conv3(x))))
        x = self.pool4(self.relu4(self.bn4(self.conv4(x))))
        x = x.view(-1, 1024 * 2 * 2)
        x = self.relu5(self.bn5(self.fc1(x)))
        x = self.fc2(x)
        return x

net = Net() # Create the network instance.
```

```
net.to(device) # Move the network parameters to the specified device.
```

```
[5]: Net(
  (conv1): Conv2d(3, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu1): ReLU()
  (conv2): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool2): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu2): ReLU()
  (conv3): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool3): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu3): ReLU()
  (conv4): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn4): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool4): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu4): ReLU()
  (fc1): Linear(in_features=4096, out_features=2048, bias=True)
  (bn5): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu5): ReLU()
  (fc2): Linear(in_features=2048, out_features=100, bias=True)
)
```

### Optimizer and Loss Function

```
[6]: # We use cross-entropy as loss function.
loss_func = nn.CrossEntropyLoss()
# We use stochastic gradient descent (SGD) as optimizer.
#opt = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)
opt = optim.Adam(net.parameters(), lr=0.0001)
```

### Training Procedure

```
[7]: import sys
from tqdm.notebook import tqdm

avg_losses = [] # Avg. losses.
epochs = 8 # Total epochs.
print_freq = 500 # Print frequency.
```

```

for epoch in range(epochs): # Loop over the dataset multiple times.
    running_loss = 0.0      # Initialize running loss.
    for i, data in enumerate(tqdm(trainloader), 0):
        # Get the inputs.
        inputs, labels = data

        # Move the inputs to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients.
        opt.zero_grad()

        # Forward step.
        outputs = net(inputs)
        loss = loss_func(outputs, labels)

        # Backward step.
        loss.backward()

        # Optimization step (update the parameters).
        opt.step()

        # Print statistics.
        running_loss += loss.item()
        if i % print_freq == print_freq - 1: # Print every several mini-batches.
            avg_loss = running_loss / print_freq
            print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.
                ↪format(epoch, i, avg_loss), flush=True)
            sys.stdout.flush()
            avg_losses.append(avg_loss)
            running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

[epoch: 0, i: 499] avg mini-batch loss: 4.601
[epoch: 0, i: 999] avg mini-batch loss: 4.316
[epoch: 0, i: 1499] avg mini-batch loss: 4.208
[epoch: 0, i: 1999] avg mini-batch loss: 4.090
[epoch: 0, i: 2499] avg mini-batch loss: 4.026
[epoch: 0, i: 2999] avg mini-batch loss: 3.944
[epoch: 0, i: 3499] avg mini-batch loss: 3.837
[epoch: 0, i: 3999] avg mini-batch loss: 3.816
[epoch: 0, i: 4499] avg mini-batch loss: 3.778
[epoch: 0, i: 4999] avg mini-batch loss: 3.765
[epoch: 0, i: 5499] avg mini-batch loss: 3.656
[epoch: 0, i: 5999] avg mini-batch loss: 3.606

```

```

[epoch: 0, i: 6499] avg mini-batch loss: 3.593
[epoch: 0, i: 6999] avg mini-batch loss: 3.497
[epoch: 0, i: 7499] avg mini-batch loss: 3.557
[epoch: 0, i: 7999] avg mini-batch loss: 3.394
[epoch: 0, i: 8499] avg mini-batch loss: 3.415
[epoch: 0, i: 8999] avg mini-batch loss: 3.317
[epoch: 0, i: 9499] avg mini-batch loss: 3.322
[epoch: 0, i: 9999] avg mini-batch loss: 3.270
[epoch: 0, i: 10499] avg mini-batch loss: 3.335
[epoch: 0, i: 10999] avg mini-batch loss: 3.158
[epoch: 0, i: 11499] avg mini-batch loss: 3.179
[epoch: 0, i: 11999] avg mini-batch loss: 3.159
[epoch: 0, i: 12499] avg mini-batch loss: 3.160

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 1, i: 499] avg mini-batch loss: 3.079
[epoch: 1, i: 999] avg mini-batch loss: 3.013
[epoch: 1, i: 1499] avg mini-batch loss: 3.008
[epoch: 1, i: 1999] avg mini-batch loss: 3.020
[epoch: 1, i: 2499] avg mini-batch loss: 2.979
[epoch: 1, i: 2999] avg mini-batch loss: 2.942
[epoch: 1, i: 3499] avg mini-batch loss: 2.924
[epoch: 1, i: 3999] avg mini-batch loss: 2.953
[epoch: 1, i: 4499] avg mini-batch loss: 2.916
[epoch: 1, i: 4999] avg mini-batch loss: 2.904
[epoch: 1, i: 5499] avg mini-batch loss: 2.908
[epoch: 1, i: 5999] avg mini-batch loss: 2.843
[epoch: 1, i: 6499] avg mini-batch loss: 2.875
[epoch: 1, i: 6999] avg mini-batch loss: 2.885
[epoch: 1, i: 7499] avg mini-batch loss: 2.887
[epoch: 1, i: 7999] avg mini-batch loss: 2.833
[epoch: 1, i: 8499] avg mini-batch loss: 2.792
[epoch: 1, i: 8999] avg mini-batch loss: 2.831
[epoch: 1, i: 9499] avg mini-batch loss: 2.738
[epoch: 1, i: 9999] avg mini-batch loss: 2.765
[epoch: 1, i: 10499] avg mini-batch loss: 2.700
[epoch: 1, i: 10999] avg mini-batch loss: 2.791
[epoch: 1, i: 11499] avg mini-batch loss: 2.665
[epoch: 1, i: 11999] avg mini-batch loss: 2.688
[epoch: 1, i: 12499] avg mini-batch loss: 2.664

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

```

```

[epoch: 2, i: 499] avg mini-batch loss: 2.572
[epoch: 2, i: 999] avg mini-batch loss: 2.482
[epoch: 2, i: 1499] avg mini-batch loss: 2.580
[epoch: 2, i: 1999] avg mini-batch loss: 2.504
[epoch: 2, i: 2499] avg mini-batch loss: 2.608

```

IOPub message rate exceeded.  
The notebook server will temporarily stop sending output  
to the client in order to avoid crashing it.  
To change this limit, set the config variable  
`--NotebookApp.iopub\_msg\_rate\_limit`.

Current values:  
NotebookApp.iopub\_msg\_rate\_limit=1000.0 (msgs/sec)  
NotebookApp.rate\_limit\_window=3.0 (secs)

```
[epoch: 4, i: 3499] avg mini-batch loss: 1.966
[epoch: 4, i: 3999] avg mini-batch loss: 2.028
[epoch: 4, i: 4499] avg mini-batch loss: 2.013
[epoch: 4, i: 4999] avg mini-batch loss: 2.017
[epoch: 4, i: 5499] avg mini-batch loss: 2.001
[epoch: 4, i: 5999] avg mini-batch loss: 2.013
[epoch: 4, i: 6499] avg mini-batch loss: 2.053
[epoch: 4, i: 6999] avg mini-batch loss: 2.057
[epoch: 4, i: 7499] avg mini-batch loss: 2.050
[epoch: 4, i: 7999] avg mini-batch loss: 1.978
[epoch: 4, i: 8499] avg mini-batch loss: 2.131
[epoch: 4, i: 8999] avg mini-batch loss: 2.037
[epoch: 4, i: 9499] avg mini-batch loss: 2.087
[epoch: 4, i: 9999] avg mini-batch loss: 2.019
[epoch: 4, i: 10499] avg mini-batch loss: 1.987
[epoch: 4, i: 10999] avg mini-batch loss: 2.074
[epoch: 4, i: 11499] avg mini-batch loss: 2.061
[epoch: 4, i: 11999] avg mini-batch loss: 2.023
[epoch: 4, i: 12499] avg mini-batch loss: 1.960
```

0% | 0/12500 [00:00<?, ?it/s]

```
[epoch: 5, i: 499] avg mini-batch loss: 1.790
[epoch: 5, i: 999] avg mini-batch loss: 1.767
[epoch: 5, i: 1499] avg mini-batch loss: 1.766
[epoch: 5, i: 1999] avg mini-batch loss: 1.874
[epoch: 5, i: 2499] avg mini-batch loss: 1.787
[epoch: 5, i: 2999] avg mini-batch loss: 1.811
[epoch: 5, i: 3499] avg mini-batch loss: 1.797
[epoch: 5, i: 3999] avg mini-batch loss: 1.843
```

IOPub message rate exceeded.  
The notebook server will temporarily stop sending output  
to the client in order to avoid crashing it.  
To change this limit, set the config variable  
`--NotebookApp.iopub\_msg\_rate\_limit`.

Current values:  
NotebookApp.iopub\_msg\_rate\_limit=1000.0 (msgs/sec)

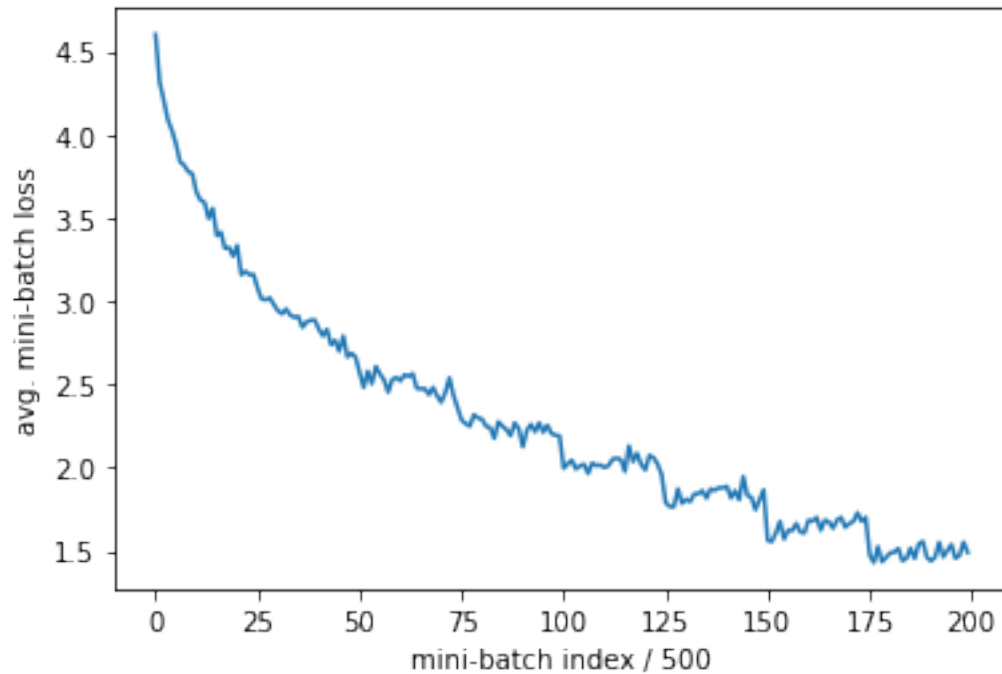
NotebookApp.rate\_limit\_window=3.0 (secs)

```
[epoch: 7, i: 2999] avg mini-batch loss: 1.489
[epoch: 7, i: 3499] avg mini-batch loss: 1.493
[epoch: 7, i: 3999] avg mini-batch loss: 1.516
[epoch: 7, i: 4499] avg mini-batch loss: 1.443
[epoch: 7, i: 4999] avg mini-batch loss: 1.456
[epoch: 7, i: 5499] avg mini-batch loss: 1.518
[epoch: 7, i: 5999] avg mini-batch loss: 1.456
[epoch: 7, i: 6499] avg mini-batch loss: 1.542
[epoch: 7, i: 6999] avg mini-batch loss: 1.558
[epoch: 7, i: 7499] avg mini-batch loss: 1.463
[epoch: 7, i: 7999] avg mini-batch loss: 1.441
[epoch: 7, i: 8499] avg mini-batch loss: 1.464
[epoch: 7, i: 8999] avg mini-batch loss: 1.553
[epoch: 7, i: 9499] avg mini-batch loss: 1.466
[epoch: 7, i: 9999] avg mini-batch loss: 1.502
[epoch: 7, i: 10499] avg mini-batch loss: 1.540
[epoch: 7, i: 10999] avg mini-batch loss: 1.457
[epoch: 7, i: 11499] avg mini-batch loss: 1.474
[epoch: 7, i: 11999] avg mini-batch loss: 1.553
[epoch: 7, i: 12499] avg mini-batch loss: 1.492
Finished Training.
```

### Training Loss Curve

```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```

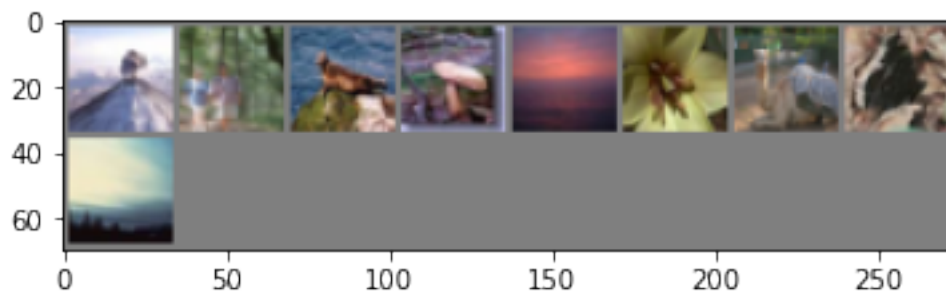




### Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                               for j in range(4)))
```



GroundTruth: mountain forest seal mushroom  
Predicted: road rabbit otter mushroom

```
[12]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 48 %

```
[ ]: # Get test accuracy for each class.
class_correct = [0] * len(classes)
class_total = [0] * len(classes)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(len(labels)):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

```
[ ]: # One of the changes I made was that I added another layer in the network
# that takes the output from the second convolutional layer, applies ReLU
↳activation,
# and passes it through a 2x2 AvgPool layer to capture more relationships.
# I also reduced the learning rate of the optimizer to half (0.0005) which
# which allowed the optimizer to take smaller steps towards the minimum of
# the loss function which might allow for a better chance of finding the global
↳min of loss.
```