

final_24

March 26, 2023

```
[1]: #!pip install torch torchvision
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

Prepare for Dataset

```
[2]: transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('apple', 'aquarium_fish', 'baby', 'bear', 'beaver',
           'bed', 'bee', 'beetle', 'bicycle', 'bottle', 'bowl',
           'boy', 'bridge', 'bus', 'butterfly', 'camel', 'can',
           'castle', 'caterpillar', 'cattle', 'chair',
           'chimpanzee', 'clock', 'cloud', 'cockroach', 'couch',
```

```
'crab', 'crocodile', 'cup', 'dinosaur', 'dolphin',
'elephant', 'flatfish', 'forest', 'fox', 'girl',
'hamster', 'house', 'kangaroo', 'keyboard', 'lamp',
'lawn_mower', 'leopard', 'lion', 'lizard', 'lobster',
'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
'mushroom', 'oak_tree', 'orange', 'orchid', 'otter',
'palm_tree', 'pear', 'pickup_truck', 'pine_tree', 'plain',
'plate', 'poppy', 'porcupine', 'possum', 'rabbit', 'raccoon',
'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
'shrew', 'skunk', 'skyscraper', 'snail', 'snake', 'spider',
'squirrel', 'streetcar', 'sunflower', 'sweet_pepper', 'table',
'tank', 'telephone', 'television', 'tiger', 'tractor', 'train',
'trout', 'tulip', 'turtle', 'wardrobe', 'whale', 'willow_tree',
'wolf', 'woman', 'worm')
```

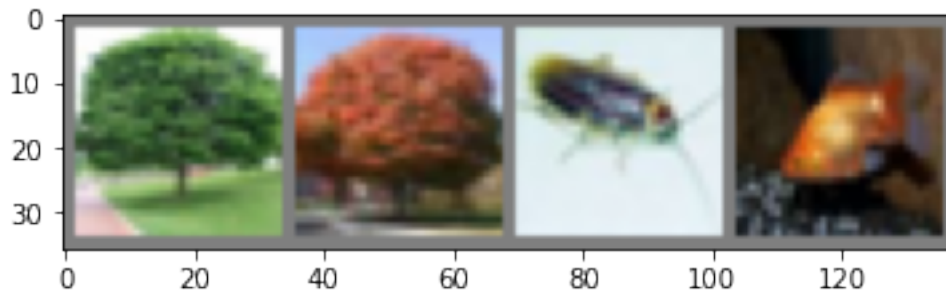
```
#classes = ('plane', 'car', 'bird', 'cat',
#           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified

```
[3]: # The function to show an image.
def imshow(img):
    img = img / 2 + 0.5     # Unnormalize.
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Get some random training images.
dataiter = iter(trainloader)
images, labels = next(dataiter)
# Show images.
imshow(torchvision.utils.make_grid(images))
# Print labels.
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



```
oak_tree maple_tree cockroach aquarium_fish
```

Choose a Device

```
[4]: # If there are GPUs, choose the first one for computing. Otherwise use CPU.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
# If 'cuda:0' is printed, it means GPU is available.
```

```
cuda:0
```

Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=128, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(128)
        self.pool1 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(128, 256, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(256)
        self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu2 = nn.ReLU()
        self.conv3 = nn.Conv2d(256, 512, 3, padding=1)
        self.bn3 = nn.BatchNorm2d(512)
        self.pool3 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu3 = nn.ReLU()
        self.conv4 = nn.Conv2d(512, 1024, 3, padding=1)
        self.bn4 = nn.BatchNorm2d(1024)
        self.pool4 = nn.AvgPool2d(kernel_size=2, stride=2)
        self.relu4 = nn.ReLU()
        self.fc1 = nn.Linear(1024 * 2 * 2, 2048)
        self.bn5 = nn.BatchNorm1d(2048)
        self.relu5 = nn.ReLU()
        self.fc2 = nn.Linear(2048, 100)

    def forward(self, x):
        x = self.pool1(self.relu1(self.bn1(self.conv1(x))))
        x = self.pool2(self.relu2(self.bn2(self.conv2(x))))
        x = self.pool3(self.relu3(self.bn3(self.conv3(x))))
        x = self.pool4(self.relu4(self.bn4(self.conv4(x))))
        #x = x.view(x.size(0), -1)
        x = x.flatten(start_dim=1)
        x = self.relu5(self.bn5(self.fc1(x)))
        x = self.fc2(x)
        return x
```

```
net = Net()      # Create the network instance.
net.to(device)  # Move the network parameters to the specified device.
```

```
[5]: Net(
  (conv1): Conv2d(3, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool1): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu1): ReLU()
  (conv2): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool2): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu2): ReLU()
  (conv3): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool3): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu3): ReLU()
  (conv4): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn4): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool4): AvgPool2d(kernel_size=2, stride=2, padding=0)
  (relu4): ReLU()
  (fc1): Linear(in_features=4096, out_features=2048, bias=True)
  (bn5): BatchNorm1d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu5): ReLU()
  (fc2): Linear(in_features=2048, out_features=100, bias=True)
)
```

Optimizer and Loss Function

```
[6]: # We use cross-entropy as loss function.
loss_func = nn.CrossEntropyLoss()
# We use stochastic gradient descent (SGD) as optimizer.
#opt = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)
opt = optim.Adam(net.parameters(), lr=0.0001)
```

Training Procedure

```
[7]: import sys
from tqdm.notebook import tqdm

avg_losses = []    # Avg. losses.
epochs = 8         # Total epochs.
print_freq = 500   # Print frequency.
```

```

for epoch in range(epochs): # Loop over the dataset multiple times.
    running_loss = 0.0 # Initialize running loss.
    for i, data in enumerate(tqdm(trainloader), 0):
        # Get the inputs.
        inputs, labels = data

        # Move the inputs to the specified device.
        inputs, labels = inputs.to(device), labels.to(device)

        # Zero the parameter gradients.
        opt.zero_grad()

        # Forward step.
        outputs = net(inputs)
        loss = loss_func(outputs, labels)

        # Backward step.
        loss.backward()

        # Optimization step (update the parameters).
        opt.step()

        # Print statistics.
        running_loss += loss.item()
        if i % print_freq == print_freq - 1: # Print every several mini-batches.
            avg_loss = running_loss / print_freq
            print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.
                ↪format(epoch, i, avg_loss), flush=True)
            sys.stdout.flush()
            avg_losses.append(avg_loss)
            running_loss = 0.0

print('Finished Training.')

```

```

0%|          | 0/12500 [00:00<?, ?it/s]

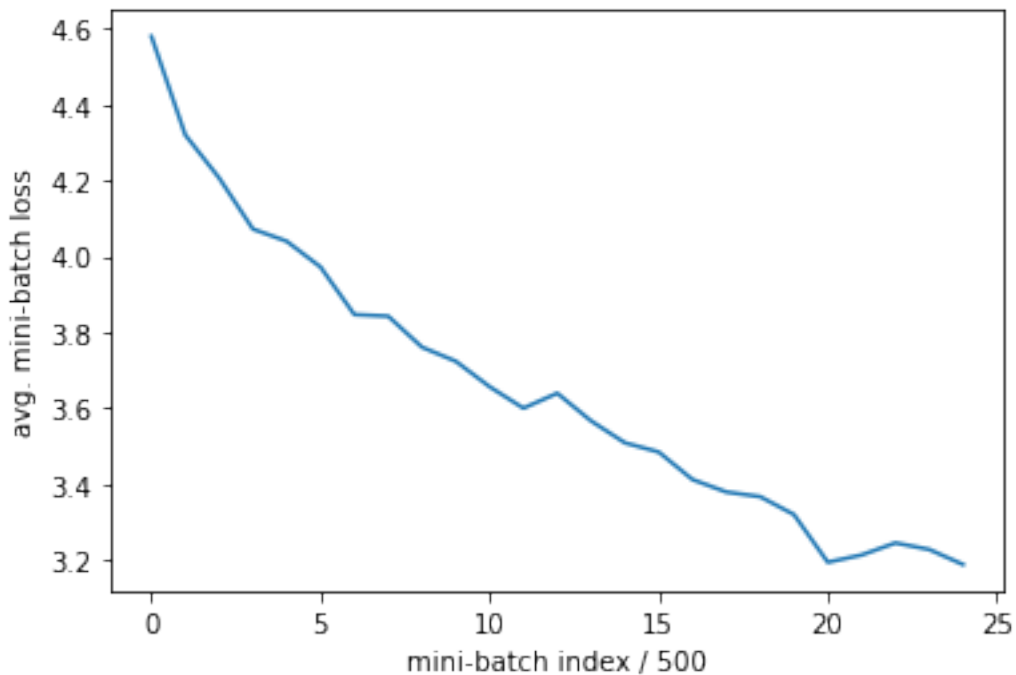
[epoch: 0, i:   499] avg mini-batch loss: 4.580
[epoch: 0, i:   999] avg mini-batch loss: 4.320
[epoch: 0, i:  1499] avg mini-batch loss: 4.207
[epoch: 0, i:  1999] avg mini-batch loss: 4.072
[epoch: 0, i:  2499] avg mini-batch loss: 4.040
[epoch: 0, i:  2999] avg mini-batch loss: 3.972
[epoch: 0, i:  3499] avg mini-batch loss: 3.846
[epoch: 0, i:  3999] avg mini-batch loss: 3.842
[epoch: 0, i:  4499] avg mini-batch loss: 3.760
[epoch: 0, i:  4999] avg mini-batch loss: 3.723
[epoch: 0, i:  5499] avg mini-batch loss: 3.656

```

```
[epoch: 0, i: 5999] avg mini-batch loss: 3.600
[epoch: 0, i: 6499] avg mini-batch loss: 3.639
[epoch: 0, i: 6999] avg mini-batch loss: 3.566
[epoch: 0, i: 7499] avg mini-batch loss: 3.508
[epoch: 0, i: 7999] avg mini-batch loss: 3.484
[epoch: 0, i: 8499] avg mini-batch loss: 3.411
[epoch: 0, i: 8999] avg mini-batch loss: 3.378
[epoch: 0, i: 9499] avg mini-batch loss: 3.366
[epoch: 0, i: 9999] avg mini-batch loss: 3.319
[epoch: 0, i: 10499] avg mini-batch loss: 3.193
[epoch: 0, i: 10999] avg mini-batch loss: 3.212
[epoch: 0, i: 11499] avg mini-batch loss: 3.244
[epoch: 0, i: 11999] avg mini-batch loss: 3.227
[epoch: 0, i: 12499] avg mini-batch loss: 3.187
Finished Training.
```

Training Loss Curve


```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                              for j in range(4)))
```



GroundTruth: mountain forest seal mushroom
Predicted: whale rabbit camel streetcar

```
[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 24 %

```
[11]: # Get test accuracy for each class.
class_correct = [0] * len(classes)
class_total = [0] * len(classes)
with torch.no_grad():
```

```

for data in testloader:
    images, labels = data
    images, labels = images.to(device), labels.to(device)
    outputs = net(images)
    _, predicted = torch.max(outputs, 1)
    c = (predicted == labels).squeeze()
    for i in range(len(labels)):
        label = labels[i]
        class_correct[label] += c[i].item()
        class_total[label] += 1

for i in range(len(classes)):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

```

Accuracy of apple : 41 %
Accuracy of aquarium_fish : 10 %
Accuracy of  baby :  4 %
Accuracy of  bear :  1 %
Accuracy of beaver : 16 %
Accuracy of  bed : 17 %
Accuracy of  bee : 16 %
Accuracy of beetle : 26 %
Accuracy of bicycle : 21 %
Accuracy of bottle : 35 %
Accuracy of  bowl : 12 %
Accuracy of  boy :  7 %
Accuracy of bridge : 16 %
Accuracy of  bus :  9 %
Accuracy of butterfly : 35 %
Accuracy of camel :  7 %
Accuracy of  can : 20 %
Accuracy of castle : 30 %
Accuracy of caterpillar : 16 %
Accuracy of cattle : 14 %
Accuracy of chair : 72 %
Accuracy of chimpanzee : 21 %
Accuracy of clock :  6 %
Accuracy of cloud : 66 %
Accuracy of cockroach : 55 %
Accuracy of couch :  1 %
Accuracy of  crab : 15 %
Accuracy of crocodile : 22 %
Accuracy of  cup : 39 %
Accuracy of dinosaur : 23 %
Accuracy of dolphin : 18 %
Accuracy of elephant : 32 %

```


Accuracy of flatfish : 13 %
Accuracy of forest : 23 %
Accuracy of fox : 3 %
Accuracy of girl : 21 %
Accuracy of hamster : 14 %
Accuracy of house : 30 %
Accuracy of kangaroo : 25 %
Accuracy of keyboard : 20 %
Accuracy of lamp : 20 %
Accuracy of lawn_mower : 54 %
Accuracy of leopard : 20 %
Accuracy of lion : 11 %
Accuracy of lizard : 8 %
Accuracy of lobster : 11 %
Accuracy of man : 10 %
Accuracy of maple_tree : 50 %
Accuracy of motorcycle : 45 %
Accuracy of mountain : 16 %
Accuracy of mouse : 4 %
Accuracy of mushroom : 10 %
Accuracy of oak_tree : 24 %
Accuracy of orange : 54 %
Accuracy of orchid : 40 %
Accuracy of otter : 6 %
Accuracy of palm_tree : 32 %
Accuracy of pear : 32 %
Accuracy of pickup_truck : 21 %
Accuracy of pine_tree : 20 %
Accuracy of plain : 55 %
Accuracy of plate : 37 %
Accuracy of poppy : 47 %
Accuracy of porcupine : 24 %
Accuracy of possum : 5 %
Accuracy of rabbit : 22 %
Accuracy of raccoon : 22 %
Accuracy of ray : 25 %
Accuracy of road : 77 %
Accuracy of rocket : 32 %
Accuracy of rose : 42 %
Accuracy of sea : 47 %
Accuracy of seal : 3 %
Accuracy of shark : 25 %
Accuracy of shrew : 14 %
Accuracy of skunk : 38 %
Accuracy of skyscraper : 61 %
Accuracy of snail : 10 %
Accuracy of snake : 0 %
Accuracy of spider : 11 %

Accuracy of squirrel : 3 %
Accuracy of streetcar : 55 %
Accuracy of sunflower : 50 %
Accuracy of sweet_pepper : 9 %
Accuracy of table : 12 %
Accuracy of tank : 30 %
Accuracy of telephone : 28 %
Accuracy of television : 42 %
Accuracy of tiger : 9 %
Accuracy of tractor : 19 %
Accuracy of train : 0 %
Accuracy of trout : 25 %
Accuracy of tulip : 4 %
Accuracy of turtle : 6 %
Accuracy of wardrobe : 45 %
Accuracy of whale : 38 %
Accuracy of willow_tree : 17 %
Accuracy of wolf : 23 %
Accuracy of woman : 17 %
Accuracy of worm : 16 %

```
[12]: # One of the changes I made was that I added another layer in the network
      # that takes the output from the second convolutional layer, applies ReLU
      ↪activation,
      # and passes it through a 2x2 AvgPool layer to capture more relationships.
      # I also reduced the learning rate of the optimizer to half (0.0005) which
      # which allowed the optimizer to take smaller steps towards the minimum of
      # the loss function which might allow for a better chance of finding the global
      ↪min of loss.
```